

Package ‘bayesmeta’

October 11, 2018

Type Package

Title Bayesian Random-Effects Meta-Analysis

Version 2.3

Date 2018-10-11

Maintainer Christian Roever <christian.roever@med.uni-goettingen.de>

Depends forestplot (>= 1.6), metafor (>= 2.0-0)

Suggests compute.es, knitr, rmarkdown, R.rsp

Description A collection of functions allowing to derive the posterior distribution of the two parameters in a random-effects meta-analysis, and providing functionality to evaluate joint and marginal posterior probability distributions, predictive distributions, shrinkage effects, posterior predictive p-values, etc.

License GPL (>= 2)

URL <https://gitlab.gwdg.de/croever/bayesmeta>,
<http://ams.med.uni-goettingen.de:3838/bayesmeta/app>

VignetteBuilder knitr, R.rsp

NeedsCompilation no

Author Christian Roever [aut, cre] (<<https://orcid.org/0000-0002-6911-698X>>),
Tim Friede [ctb] (<<https://orcid.org/0000-0001-5347-7441>>)

Repository CRAN

Date/Publication 2018-10-11 12:40:02 UTC

R topics documented:

bayesmeta-package	2
bayesmeta	3
Cochran1954	14
CrinsEtAl2014	16
dhalflogistic	18
dhalfnormal	19
dlomax	21

drayleigh	23
forest.bayesmeta	25
forestplot.bayesmeta	27
GoralczykEtAl2011	29
HinksEtAl2010	31
normalmixture	33
Peto1980	37
plot.bayesmeta	38
pppvalue	41
RhodesEtAlPrior	45
Rubin1981	48
SidikJonkman2007	50
SnedecorCochran	52
TurnerEtAlPrior	53
Index	57

bayesmeta-package	<i>Bayesian Random-Effects Meta-Analysis</i>
-------------------	--

Description

Description: A collection of functions allowing to derive the posterior distribution of the two parameters in a random-effects meta-analysis, and providing functionality to evaluate joint and marginal posterior probability distributions, predictive distributions, shrinkage effects, posterior predictive p-values, etc.

Details

Package: bayesmeta
Type: Package
Version: 2.3
Date: 2018-10-11
License: GPL (>=2)

The main functionality is provided by the [bayesmeta\(\)](#) function. It takes the data (estimates and associated standard errors) and prior information (effect and heterogeneity priors), and returns an object containing functions that allow to derive posterior quantities like joint or marginal densities, quantiles, etc.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. *arXiv preprint 1711.08683* (submitted for publication), 2017.

C. Roever. *The bayesmeta app*. <http://ams.med.uni-goettingen.de:3838/bayesmeta/app>, 2018.

See Also

[forestplot.bayesmeta](#), [plot.bayesmeta](#).

Examples

```
# example data by Snedecor and Cochran:
data("SnedecorCochran")

## Not run:
# analysis using improper uniform prior
# (may take a few seconds to compute!):
bma <- bayesmeta(y=SnedecorCochran[, "mean"],
                 sigma=sqrt(SnedecorCochran[, "var"]),
                 label=SnedecorCochran[, "no"])

# show some summary statistics:
bma

# show a bit more details:
summary(bma)

# show a forest plot:
forestplot(bma)

# show some more plots:
plot(bma)

## End(Not run)
```

bayesmeta

Bayesian random-effects meta-analysis

Description

This function allows to derive the posterior distribution of the two parameters in a random-effects meta-analysis and provides functions to evaluate joint and marginal posterior probability distributions, etc.

Usage

```
bayesmeta(y, ...)
## Default S3 method:
bayesmeta(y, sigma, labels = names(y),
          tau.prior = "uniform",
          mu.prior = c("mean"=NA, "sd"=NA),
```

```

mu.prior.mean = mu.prior[1], mu.prior.sd = mu.prior[2],
interval.type = c("shortest", "central"),
delta = 0.01, epsilon = 0.0001,
rel.tol.integrate = 2^16*.Machine$double.eps,
abs.tol.integrate = 0.0,
tol.uniroot = rel.tol.integrate, ...)
## S3 method for class 'escalc'
bayesmeta(y, labels = NULL, ...)

```

Arguments

<code>y</code>	vector of estimates <i>or</i> an escalc object.
<code>sigma</code>	vector of standard errors associated with <code>y</code> .
<code>labels</code>	(optional) a vector of labels corresponding to <code>y</code> and <code>sigma</code> .
<code>tau.prior</code>	<p>a function returning the prior density for the heterogeneity parameter (τ) <i>or</i> a character string specifying one of the <i>default 'non-informative' priors</i>; possible choices for the latter case are:</p> <ul style="list-style-type: none"> • "uniform": a uniform prior in τ • "sqrt": a uniform prior in $\sqrt{\tau}$ • "Jeffreys": the Jeffreys prior for τ • "BergerDeely": the prior due to Berger and Deely (1988) • "conventional": the conventional prior • "DuMouchel": the DuMouchel prior • "shrinkage": the 'uniform shrinkage' prior • "I2": a uniform prior on the 'relative heterogeneity' I^2 <p>The default is "uniform" (which should be used with caution; see remarks below). The above priors are described in some more detail below.</p>
<code>mu.prior</code>	the mean and standard deviation of the normal prior distribution for the effect μ . If unspecified, an (improper) uniform prior is used.
<code>mu.prior.mean</code> , <code>mu.prior.sd</code>	alternative parameters to specify the prior distribution for the effect μ (see above).
<code>interval.type</code>	the type of (credible, prediction, shrinkage) interval to be returned by default; either "shortest" for shortest intervals, or "central" for central, equal-tailed intervals.
<code>delta</code> , <code>epsilon</code>	the parameters specifying the desired accuracy for approximation of the μ posterior(s), and with that determining the number of τ support points being used internally. Smaller values imply greater accuracy and greater computational burden (Roever and Friede, 2017).
<code>rel.tol.integrate</code> , <code>abs.tol.integrate</code> , <code>tol.uniroot</code>	the <code>rel.tol</code> , <code>abs.tol</code> and <code>tol</code> 'accuracy' arguments that are passed to the integrate() or uniroot() functions for internal numerical integration or root finding (see also the help there).
<code>...</code>	other bayesmeta arguments.

Details

The common random-effects meta-analysis model may be stated as

$$y_i | \mu, \sigma_i, \tau \sim \text{Normal}(\mu, \sigma_i^2 + \tau^2)$$

where the data (y, σ) enter as y_i , the i -th estimate, that is associated with standard error $\sigma_i > 0$, and $i = 1, \dots, k$. The model includes two unknown parameters, namely the (mean) effect μ , and the heterogeneity τ . Alternatively, the model may also be formulated via an intermediate step as

$$y_i | \theta_i, \sigma_i \sim \text{Normal}(\theta_i, \sigma_i^2),$$

$$\theta_i | \mu, \tau \sim \text{Normal}(\mu, \tau^2),$$

where the θ_i denote the *trial-specific* means that are then measured through the estimate y_i with an associated measurement uncertainty σ_i . The θ_i again differ from trial to trial and are distributed around a common mean μ with standard deviation τ .

The `bayesmeta()` function utilizes the fact that the joint posterior distribution $p(\mu, \tau | y, \sigma)$ may be partly analytically integrated out to determine the integrals necessary for coherent Bayesian inference on one or both of the parameters.

As long as we assume that the prior probability distribution may be factored into independent marginals $p(\mu, \tau) = p(\mu) \times p(\tau)$ and either an (improper) uniform or a normal prior is used for the effect μ , the joint likelihood $p(y | \mu, \tau)$ may be analytically marginalized over μ , yielding the marginal likelihood function $p(y | \tau)$. Using any prior $p(\tau)$ for the heterogeneity, the 1-dimensional marginal posterior $p(\tau | y) \propto p(y | \tau) \times p(\tau)$ may then be treated numerically. As the *conditional* posterior $p(\mu | \tau, y)$ is a normal distribution, inference on the remaining joint $(p(\mu, \tau | y))$ or marginal $(p(\mu | y))$ posterior may be approached numerically (using the `DIRECT` algorithm) as well. Accuracy of the computations is determined by the `delta` (maximum divergence δ) and `epsilon` (tail probability ϵ) parameters (Roeve and Friede, 2017).

What constitutes a sensible prior for both μ and τ depends (as usual) very much on the context. Potential candidates for informative (or weakly informative) heterogeneity (τ) priors may include the *half-normal*, *half-Student-t*, *half-Cauchy*, *exponential*, or *Lomax* distributions; we recommend the use of heavy-tailed prior distributions if in case of prior/data conflict the prior is supposed to be discounted (O'Hagan and Pericchi, 2012). A sensible informative prior might also be a *log-normal* or a *scaled inverse χ^2* distribution. One might argue that an uninformative prior for τ may be uniform or monotonically decreasing in τ . Another option is to use the *Jeffreys prior* (see the `tau.prior` argument above). The Jeffreys prior implemented here is the variant also described by Tibshirani (1989) that results from fixing the location parameter (μ) and considering the Fisher information element corresponding to the heterogeneity τ only. This prior also constitutes the *Berger/Bernardo reference prior* for the present problem (Bodnar et al., 2016). The *uniform shrinkage* and *Du-Mouchel* priors are described in Spiegelhalter et al. (2004, Sec. 5.7.3). The procedure is able to handle improper priors (and does so by default), but of course the usual care must be taken here, as the resulting posterior *may or may not* be a proper probability distribution.

Note that a wide range of different types of endpoints may be treated on a continuous scale after an appropriate transformation; for example, count data may be handled by considering corresponding logarithmic odds ratios. Many such transformations are implemented in the **metafor** package's `escalc` function and may be directly used as an input to the `bayesmeta()` function; see also the example below. Alternatively, the **compute.es** package also provides a range of effect sizes to be computed from different data types.

The `bayesmeta()` function eventually generates some basic summary statistics, but most importantly it provides an object containing a range of functions allowing to evaluate posterior distributions; this includes joint and marginal posteriors, prior and likelihood, predictive distributions, densities, cumulative distributions and quantile functions. For more details see also the documentation and examples below. Use of the `individual` argument allows to access posteriors of study-specific (*shrinkage*-) effects (θ_i). The `predict` argument may be used to access the predictive distribution of a future study's effect (θ_{k+1}), facilitating a *meta-analytic-predictive (MAP)* approach (Neuenschwander et al., 2010).

Prior specification details: When specifying the `tau.prior` argument as a character string (and not as a prior density function), then the actual prior probability density functions corresponding to the possible choices of the `tau.prior` argument are given by:

- "uniform" - the (improper) uniform prior in τ :

$$p(\tau) \propto 1$$

- "sqrt" - the (improper) uniform prior in $\sqrt{\tau}$:

$$p(\tau) \propto \tau^{-1/2} = \frac{1}{\sqrt{\tau}}$$

- "Jeffreys" - *Tibshirani's noninformative prior*, a variation of the *Jeffreys prior*, which here also constitutes the *Berger/Bernardo reference prior* for τ :

$$p(\tau) \propto \sqrt{\sum_{i=1}^k \left(\frac{\tau}{\sigma_i^2 + \tau^2} \right)^2}$$

(This is also an improper prior whose density does not integrate to 1).

- "BergerDeely" - the (improper) *Berger/Deely* prior:

$$p(\tau) \propto \prod_{i=1}^k \left(\frac{\tau}{\sigma_i^2 + \tau^2} \right)^{1/k}$$

This is a variation of the above *Jeffreys* prior, and both are equal in case all standard errors (σ_i) are the same.

- "conventional" - the (proper) *conventional prior*:

$$p(\tau) \propto \prod_{i=1}^k \left(\frac{\tau}{(\sigma_i^2 + \tau^2)^{3/2}} \right)^{1/k}$$

This is a proper variation of the above *Berger/Deely* prior intended especially for testing and model selection purposes.

- "DuMouchel" - the (proper) *DuMouchel* prior:

$$p(\tau) = \frac{s_0}{(s_0 + \tau)^2}$$

where $s_0 = \sqrt{s_0^2}$ and s_0^2 again is the harmonic mean of the standard errors (as above).

- "shrinkage" - the (proper) *uniform shrinkage* prior:

$$p(\tau) = \frac{2s_0^2\tau}{(s_0^2 + \tau^2)^2}$$

where $s_0^2 = \frac{k}{\sum_{i=1}^k \sigma_i^{-2}}$ is the harmonic mean of the squared standard errors σ_i^2 .

- "I2" - the (proper) uniform prior in I^2 :

$$p(\tau) = \frac{2\hat{\sigma}^2\tau}{(\hat{\sigma}^2 + \tau^2)^2}$$

where $\hat{\sigma}^2 = \frac{(k-1) \sum_{i=1}^k \sigma_i^{-2}}{(\sum_{i=1}^k \sigma_i^{-2})^2 - \sum_{i=1}^k \sigma_i^{-4}}$. This prior is similar to the uniform shrinkage prior, except for the use of $\hat{\sigma}^2$ instead of s_0^2 .

For empirically motivated informative heterogeneity priors see also the [TurnerEtAlPrior\(\)](#) and [RhodesEtAlPrior\(\)](#) functions.

Credible intervals: Credible intervals (as well as prediction and shrinkage intervals) may be determined in different ways. By default, *shortest* intervals are returned, which for unimodal posteriors (the usual case) is equivalent to the *highest posterior density region* (Gelman et al., 1997, Sec. 2.3). Alternatively, central (equal-tailed) intervals may also be derived. The default behaviour may be controlled via the `interval.type` argument, or also by using the `method` argument with each individual call of the `$post.interval()` function (see below). A third option, although not available for prediction or shrinkage intervals, and hence not as an overall default choice, but only for the `$post.interval()` function, is to determine the *evidentiary* credible interval, which has the advantage of being parameterization invariant (Shalloway, 2014).

Bayes factors: Bayes factors (Kass and Raftery, 1995) for the two hypotheses of $\tau = 0$ and $\mu = 0$ are provided in the `$bayesfactor` element; *low* or *high* values here constitute evidence *against* or *in favour of* the hypotheses, respectively. Bayes factors are based on marginal likelihoods and can only be computed if the priors for heterogeneity and effect are proper. Bayes factors for other hypotheses can be computed using the marginal likelihood (as provided through the `$marginal` element) and the `$likelihood()` function. Bayes factors must be interpreted with very much caution, as they are susceptible to *Lindley's paradox* (Lindley, 1957), which especially implies that variations of the prior specifications that have only minuscule effects on the posterior distribution may have a substantial impact on Bayes factors (via the marginal likelihood). For more details on the problems and challenges related to Bayes factors see also Gelman et al. (1997, Sec. 7.4).

Besides the 'actual' Bayes factors, *minimum Bayes factors* are also provided (Spiegelhalter et al., 2004; Sec. 4.4). The minimum Bayes factor for the hypothesis of $\mu = 0$ constitutes the minimum across all possible priors for μ and hence gives a measure of how much (or how little) evidence *against* the hypothesis is provided by the data *at most*. It is independent of the particular effect prior used in the analysis, but still dependent on the heterogeneity prior. Analogously, the same is true for the heterogeneity's minimum Bayes factor. A minimum Bayes factor can also be computed when only one of the priors is proper.

Numerical accuracy: Accuracy of the numerical results is determined by four parameters, namely, the accuracy of numerical integration as specified through the `rel.tol.integrate` and `abs.tol.integrate` arguments (which are internally passed on to the [integrate](#) function), and the accuracy of the grid approximation used for integrating out the heterogeneity as specified

through the `delta` and `epsilon` arguments (Roever and Friede, 2017). As these may also heavily impact on the computation time, be careful when changing these from their default values. You can monitor the effect of different settings by checking the number and range of support points returned in the `$support` element.

Study weights: Conditional on a given τ value, the posterior expectations of the overall effect (μ) as well as the shrinkage estimates (θ_i) result as convex combinations of the estimates y_i . The *weights* associated with each estimate y_i are commonly quoted in frequentist meta-analysis results in order to quantify (arguably somewhat heuristically) each study's contribution to the overall estimates, often in terms of percentages.

In a Bayesian meta-analysis, these numbers do not immediately arise, since the heterogeneity is marginalized over. However, due to linearity, the posterior mean effects may still be expressed in terms of linear combinations of initial estimates y_i , with weights now given by the *posterior mean weights*, marginalized over the heterogeneity τ . The posterior mean weights are returned in the `$weights` and `$weights.theta` elements, for the overall effect μ as well as for the shrinkage estimates θ_i .

Value

A list of class `bayesmeta` containing the following elements:

<code>y</code>	vector of estimates (the input data).
<code>sigma</code>	vector of standard errors corresponding to <code>y</code> (input data).
<code>labels</code>	vector of labels corresponding to <code>y</code> and <code>sigma</code> .
<code>k</code>	number of data points (in <code>y</code>).
<code>tau.prior</code>	the prior probability density function for τ .
<code>mu.prior.mean</code>	the prior mean of μ .
<code>mu.prior.sd</code>	the prior standard deviation of μ .
<code>dprior</code>	a function(<code>tau=NA</code> , <code>mu=NA</code> , <code>log=FALSE</code>) of two parameters, <code>tau</code> and/or <code>mu</code> , returning either the joint or marginal prior probability density, depending on which parameter(s) is/are provided.
<code>tau.prior.proper</code>	a logical flag indicating whether the heterogeneity prior appears to be proper (which is judged based on an attempted numerical integration of the density function).
<code>likelihood</code>	a function(<code>tau=NA</code> , <code>mu=NA</code> , <code>log=FALSE</code>) of two parameters, <code>tau</code> and/or <code>mu</code> , returning either the joint or marginal likelihood, depending on which parameter(s) is/are provided.
<code>dposterior</code>	a function(<code>tau=NA</code> , <code>mu=NA</code> , <code>theta=mu</code> , <code>log=FALSE</code> , <code>predict=FALSE</code> , <code>individual=FALSE</code>) of two parameters, <code>tau</code> and/or <code>mu</code> , returning either the joint or marginal posterior probability density, depending on which parameter(s) is/are provided. Using the argument <code>predict=TRUE</code> yields the <i>posterior predictive distribution</i> for θ . Using the <code>individual</code> argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label.

pposterior	a function(τ =NA, μ =NA, θ =mu, predict=FALSE, individual=FALSE) of one parameter (either τ or μ) returning the corresponding marginal posterior cumulative distribution function. Using the argument predict=TRUE yields the posterior predictive distribution for θ . Using the individual argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label.
qposterior	a function(τ .p=NA, μ .p=NA, θ .p=mu.p, predict=FALSE, individual=FALSE) of one parameter (either τ .p or μ .p) returning the corresponding marginal posterior quantile function. Using the argument predict=TRUE yields the posterior predictive distribution for θ . Using the individual argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label.
rposterior	a function(n=1, predict=FALSE, individual=FALSE, τ .sample=TRUE) generating n independent random draws from the (2-dimensional) posterior distribution. Using the argument predict=TRUE yields the posterior predictive distribution for θ . Using the individual argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label. In general, this via the inversion method, so it is rather slow. However, if one is not interested in sampling the heterogeneity parameter (τ), using ' τ .sample=FALSE' will speed up the function substantially.
post.interval	a function(τ .level=NA, μ .level=NA, θ .level=mu.level, method=c("shortest", "central"), predict=FALSE, individual=FALSE) returning a credible interval, depending on which of the two parameters is provided (either τ .level or μ .level). The additional parameter method may be used to specify the desired type of interval: method = "shortest" returns the shortest interval, method = "central" returns a central interval, and method = "evidentiary" returns an evidentiary interval (Shalloway, 2014); the former is the default option. Using the argument predict=TRUE yields a posterior predictive interval for θ . Using the individual argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label.
cond.moment	a function(τ , predict=FALSE, individual=FALSE, simplify=TRUE) returning conditional moments (mean and standard deviation) of μ as a function of τ . Using the argument predict=TRUE yields (conditional) posterior predictive moments for θ . Using the individual argument, you can request individual effects' (θ_i) posterior distributions. May be an integer number (1, ..., k) giving the index, or a character string giving the label.
I2	a function(τ) returning the 'relative' heterogeneity I^2 as a function of τ .
summary	a matrix listing some summary statistics, namely marginal posterior mode, median, mean, standard deviation and a (shortest) 95% credible intervals, of the marginal posterior distributions of τ and μ , and of the posterior predictive distribution of θ .
interval.type	the interval.type input argument specifying the type of interval to be returned by default.
ML	a matrix giving joint and marginal maximum-likelihood estimates of (τ , μ).
MAP	a matrix giving joint and marginal maximum-a-posteriori estimates of (τ , μ).

<code>theta</code>	a matrix giving the ‘shrinkage estimates’, i.e. summary statistics of the trial-specific means θ_i .
<code>weights</code>	a vector giving the posterior expected <i>inverse-variance weights</i> for each study (and for the effect prior mean, if the effect prior was proper).
<code>weights.theta</code>	a matrix whose columns give the posterior expected weights of each study (and of the effect prior mean, if the effect prior was proper) for all shrinkage estimates.
<code>marginal.likelihood</code>	the marginal likelihood of the data (this number is only computed if proper effect and heterogeneity priors are specified).
<code>bayesfactor</code>	Bayes factors and minimum Bayes factors for the two hypotheses of $\tau = 0$ and $\mu = 0$. These depend on the marginal likelihood and hence can only be computed if proper effect and/or heterogeneity priors are specified; see also remark above.
<code>support</code>	a matrix giving the τ support points used internally in the grid approximation, along with their associated weights, conditional mean and standard deviation of μ , and the standard deviation of the (conditional) predictive distribution of θ .
<code>delta, epsilon</code>	the ‘delta’ and ‘epsilon’ input parameter determining numerical accuracy.
<code>rel.tol.integrate, abs.tol.integrate, tol.uniroot</code>	the input parameters determining the numerical accuracy of the internally used <code>integrate()</code> and <code>uniroot()</code> functions.
<code>call</code>	an object of class <code>call</code> giving the function call that generated the bayesmeta object.
<code>init.time</code>	the computation time (in seconds) used to generate the bayesmeta object.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. *arXiv preprint 1711.08683* (submitted for publication), 2017.
- C. Roever, T. Friede. *Discrete approximation of a mixture distribution via restricted divergence*. *Journal of Computational and Graphical Statistics*, **26**(1):217-222, 2017.
- J.O. Berger, J. Deely. *A Bayesian approach to ranking and selection of related means with alternatives to analysis-of-variance methodology*. *Journal of the American Statistical Association*, **83**(402):364-373, 1988.
- O. Bodnar, A. Link, C. Elster. *Objective Bayesian inference for a generalized marginal random effects model*. *Bayesian Analysis*, **11**(1):25-45, 2016.
- A. Gelman, J.B. Carlin, H.S. Stern, D.B. Rubin. *Bayesian data analysis*. Chapman & Hall / CRC, Boca Raton, 1997.
- A. Gelman. Prior distributions for variance parameters in hierarchical models. *Bayesian Analysis*, **1**(3):515-534, 2006.

- J. Hartung, G. Knapp, B.K. Sinha. *Statistical meta-analysis with applications*. Wiley, Hoboken, 2008.
- L.V. Hedges, I. Olkin. *Statistical methods for meta-analysis*. Academic Press, San Diego, 1985
- A.E. Kass, R.E. Raftery. **Bayes factors**. *Journal of the American Statistical Association*, **90**(430):773-795, 1995.
- D.V. Lindley. **A statistical paradox**. *Biometrika*, **44**(1/2):187-192, 1957.
- B. Neuenschwander, G. Capkun-Niggli, M. Branson, D.J. Spiegelhalter. **Summarizing historical information on controls in clinical trials**. *Trials*. **7**(1):5-18, 2010.
- A. O'Hagan, L. Pericchi. **Bayesian heavy-tailed models and conflict resolution: A review**. *Brazilian Journal of Probability and Statistics*. **26**(4):372-401, 2012.
- S. Shalloway. **The evidentiary credible region**. *Bayesian Analysis*, **9**(4):909-922, 2014.
- D.J. Spiegelhalter, K.R. Abrams, J.P. Myles. *Bayesian approaches to clinical trials and health-care evaluation*. Wiley & Sons, 2004.
- R. Tibshirani. **Noninformative priors for one parameter of many**. *Biometrika*, **76**(3):604-608, 1989.
- W. Viechtbauer. **Conducting meta-analyses in R with the metafor package**. *Journal of Statistical Software*, **36**(3):1-48, 2010.

See Also

[forestplot.bayesmeta](#), [plot.bayesmeta](#), [escalc](#), [compute.es](#).

Examples

```
#####
# example data by Snedecor and Cochran:
data("SnedecorCochran")

## Not run:
# analysis using improper uniform prior
# (may take a few seconds to compute!):
ma01 <- bayesmeta(y=SnedecorCochran[, "mean"], sigma=sqrt(SnedecorCochran[, "var"]),
                  label=SnedecorCochran[, "no"])

# analysis using an informative prior
# (normal for mu and half-Cauchy for tau (scale=10))
# (may take a few seconds to compute!):
ma02 <- bayesmeta(y=SnedecorCochran[, "mean"], sigma=sqrt(SnedecorCochran[, "var"]),
                  label=SnedecorCochran[, "no"],
                  mu.prior.mean=50, mu.prior.sd=50,
                  tau.prior=function(x){return(dhalfcauchy(x, scale=10))})

# show some summary statistics:
print(ma01)
summary(ma01)

# show some plots:
forestplot(ma01)
plot(ma01)
```

```

# compare resulting marginal densities;
# the effect parameter (mu):
mu <- seq(30, 80, le=200)
plot(mu, ma02$dposterior(mu=mu), type="l", lty="dashed",
      xlab=expression("effect " * mu),
      ylab=expression("marginal posterior density"),
      main="Snedecor/Cochran example")
lines(mu, ma01$dposterior(mu=mu), lty="solid")

# the heterogeneity parameter (tau):
tau <- seq(0, 50, le=200)
plot(tau, ma02$dposterior(tau=tau), type="l", lty="dashed",
      xlab=expression("heterogeneity " * tau),
      ylab=expression("marginal posterior density"),
      main="Snedecor/Cochran example")
lines(tau, ma01$dposterior(tau=tau), lty="solid")

# compute posterior median relative heterogeneity I-squared:
ma01$I2(tau=ma01$summary["median", "tau"])

# determine 90 percent upper limits on the heterogeneity tau:
ma01$qposterior(tau=0.90)
ma02$qposterior(tau=0.90)
# determine shortest 90 percent credible interval for tau:
ma01$post.interval(tau.level=0.9, method="shortest")
## End(Not run)

#####
# example data by Sidik and Jonkman:
data("SidikJonkman2007")
# add log-odds-ratios and corresponding standard errors:
sj <- SidikJonkman2007
sj <- cbind(sj, "log.or"=log(sj[, "lihr.events"])-log(sj[, "lihr.cases"]-sj[, "lihr.events"])
            -log(sj[, "oihr.events"])+log(sj[, "oihr.cases"]-sj[, "oihr.events"]),
            "log.or.se"=sqrt(1/sj[, "lihr.events"] + 1/(sj[, "lihr.cases"]-sj[, "lihr.events"])
            + 1/sj[, "oihr.events"] + 1/(sj[, "oihr.cases"]-sj[, "oihr.events"])))

## Not run:
# analysis using weakly informative half-normal prior
# (may take a few seconds to compute!):
ma03a <- bayesmeta(y=sj[, "log.or"], sigma=sj[, "log.or.se"],
                  label=sj[, "id.sj"],
                  tau.prior=function(t){dhalfnormal(t, scale=1)})

# alternatively: may utilize "metafor" package's "escalc()" function
# to compute log-ORs and standard errors:
require("metafor")
es <- escalc(measure="OR",
             ai=lihr.events, n1i=lihr.cases,
             ci=oihr.events, n2i=oihr.cases,
             slab=id, data=SidikJonkman2007)

```

```

# apply "bayesmeta()" function directly to "escalc" object:
ma03b <- bayesmeta(es, tau.prior=function(t){dhalfnormal(t,scale=1)})
# "ma03a" and "ma03b" should be identical:
print(ma03a$summary)
print(ma03b$summary)
# compare to metafor's (frequentist) random-effects meta-analysis:
rma03a <- rma.uni(es)
rma03b <- rma.uni(es, method="EB", knha=TRUE)
# compare mu estimates (estimate and confidence interval):
plot(ma03b, which=3)
abline(v=c(rma03a$b, rma03a$ci.lb, rma03a$ci.ub), col="red", lty=c(1,2,2))
abline(v=c(rma03b$b, rma03b$ci.lb, rma03b$ci.ub), col="green3", lty=c(1,2,2))
# compare tau estimates (estimate and confidence interval):
plot(ma03b, which=4)
abline(v=confint(rma03a)$random["tau",], col="red", lty=c(1,2,2))
abline(v=confint(rma03b)$random["tau",], col="green3", lty=c(1,3,3))

# show heterogeneity's posterior density:
plot(ma03a, which=4, main="Sidik/Jonkman example")

# show some numbers (mode, median and mean):
abline(v=ma03a$summary[c("mode","median","mean"),"tau"], col="blue")

# compare with Sidik and Jonkman's estimates:
sj.estimates <- sqrt(c("MM" = 0.429, # method of moments estimator
                      "VC" = 0.841, # variance component type estimator
                      "ML" = 0.562, # maximum likelihood estimator
                      "REML" = 0.598, # restricted maximum likelihood estimator
                      "EB" = 0.703, # empirical Bayes estimator
                      "MV" = 0.818, # model error variance estimator
                      "MVvc" = 0.747)) # a variation of the MV estimator
abline(v=sj.estimates, col="red", lty="dashed")
## End(Not run)

#####
# example data by Cochran:
data("Cochran1954")

## Not run:
# analysis using improper uniform prior
# (may take a few seconds to compute!):
ma04 <- bayesmeta(y=Cochran1954[, "mean"], sigma=sqrt(Cochran1954[, "se2"]),
                  label=Cochran1954[, "observer"])

# show joint posterior density:
plot(ma04, which=2, main="Cochran example")
# show (known) true parameter value:
abline(h=161)

# pick a point estimate for tau:
tau <- ma04$summary["median","tau"]
# highlight two point hypotheses (fixed vs. random effects):

```

```

abline(v=c(0, tau), col="orange", lty="dotted", lwd=2)

# show marginal posterior density:
plot(ma04, which=3)
abline(v=161)
# show the conditional distributions of the effect mu
# at two tau values corresponding to fixed and random effects models:
cm <- ma04$cond.moment(tau=c(0,tau))
mu <- seq(130,200, le=200)
lines(mu, dnorm(mu, mean=cm[1,"mean"], sd=cm[1,"sd"]), col="orange", lwd=2)
lines(mu, dnorm(mu, mean=cm[2,"mean"], sd=cm[2,"sd"]), col="orange", lwd=2)

# determine a range of tau values:
tau <- seq(0, ma04$qposterior(tau=0.99), length=100)
# compute conditional posterior moments:
cm.overall <- ma04$cond.moment(tau=tau)
# compute study-specific conditional posterior moments:
cm.indiv <- ma04$cond.moment(tau=tau, individual=TRUE)
# show forest plot along with conditional posterior means:
par(mfrow=c(1,2))
plot(ma04, which=1, main="Cochran 1954 example")
matplot(tau, cm.indiv[, "mean", ], type="l", lty="solid", col=1:ma04$k,
        xlim=c(0,max(tau)*1.2), xlab=expression("heterogeneity "*tau),
        ylab=expression("(conditional) shrinkage estimate E["*
                        theta[i]*"|"*list(tau, y, sigma)*"]"))
text(rep(max(tau)*1.01, ma04$k), cm.indiv[length(tau), "mean", ],
     ma04$label, col=1:ma04$k, adj=c(0,0.5))
lines(tau, cm.overall[, "mean", ], lty="dashed", lwd=2)
text(max(tau)*1.01, cm.overall[length(tau), "mean", ],
     "overall", adj=c(0,0.5))
par(mfrow=c(1,1))

# show the individual effects' posterior distributions:
theta <- seq(120, 240, le=300)
plot(range(theta), c(0,0.1), type="n", xlab=expression(theta[i]), ylab="")
for (i in 1:ma04$k) {
  # draw estimate +/- uncertainty as a Gaussian:
  lines(theta, dnorm(theta, mean=ma04$y[i], sd=ma04$sigma[i]), col=i+1, lty="dotted")
  # draw effect's posterior distribution:
  lines(theta, ma04$dposterior(theta=theta, indiv=i), col=i+1, lty="solid")
}
abline(h=0)
legend(max(theta), 0.1, legend=ma04$label, col=(1:ma04$k)+1, pch=15, xjust=1, yjust=1)

## End(Not run)

```

Description

This data set gives average estimated counts of flies along with standard errors from 7 different observers.

Usage

```
data("Cochran1954")
```

Format

The data frame contains the following columns:

observer	character	identifier
mean	numeric	mean count
se2	numeric	<i>squared</i> standard error

Details

Quoting from Cochran (1954), example 3, p.119: “In studies by the U.S. Public Health Service of observers’ abilities to count the number of flies which settle momentarily on a grill, each of 7 observers was shown, for a brief period, grills with known numbers of flies impaled on them and asked to estimate the numbers. For a given grill, each observer made 5 independent estimates. The data in table 9 are for a grill which actually contained 161 flies. Estimated variances are based on 4 degrees of freedom each. [...] The only point of interest in estimating the overall mean is to test whether there is any consistent bias among observers in estimating the 161 flies on the grill. Although inspection of table 9 suggests no such bias, the data will serve to illustrate the application of partial weighting.”

Source

W.G. Cochran. The combination of estimates from different experiments. *Biometrics*, **10**(1):101-129, 1954.

Examples

```
data("Cochran1954")
## Not run:
# analysis using improper uniform prior
# (may take a few seconds to compute!):
bma <- bayesmeta(y=Cochran1954[, "mean"], sigma=sqrt(Cochran1954[, "se2"]),
  label=Cochran1954[, "observer"])

# show joint posterior density:
plot(bma, which=2, main="Cochran example")
# show (known) true parameter value:
abline(h=161)

# show forest plot:
forestplot(bma, zero=161)

## End(Not run)
```

CrinsEtAl2014

*Pediatric liver transplant example data***Description**

Numbers of cases (transplant patients) and events (acute rejections, steroid resistant rejections, PTLDs, and deaths) in experimental and control groups of six studies.

Usage

```
data("CrinsEtAl2014")
```

Format

The data frame contains the following columns:

publication	character	publication identifier (first author and publication year)
year	numeric	publication year
randomized	factor	randomization status (y/n)
control.type	factor	type of control group ('concurrent' or 'historical')
comparison	factor	type of comparison ('IL-2RA only', 'delayed CNI', or 'no/low steroids')
IL2RA	factor	type of interleukin-2 receptor antagonist (IL-2RA) ('basiliximab' or 'daclizumab')
CNI	factor	type of calcineurin inhibitor (CNI) ('tacrolimus' or 'cyclosporine A')
MMF	factor	use of mycophenolate mofetil (MMF) (y/n)
followup	numeric	follow-up time in months
treat.AR.events	numeric	number of AR events in experimental group
treat.SRR.events	numeric	number of SRR events in experimental group
treat.PTLD.events	numeric	number of PTLD events in experimental group
treat.deaths	numeric	number of deaths in experimental group
treat.total	numeric	number of cases in experimental group
control.AR.events	numeric	number of AR events in control group
control.SRR.events	numeric	number of SRR events in control group
control.PTLD.events	numeric	number of PTLD events in control group
control.deaths	numeric	number of deaths in control group
control.total	numeric	number of cases in control group

Details

A systematic literature review investigated the evidence on the effect of Interleukin-2 receptor antagonists (IL-2RA) and resulted in six controlled studies reporting acute rejection (AR), steroid-resistant rejection (SRR) and post-transplant lymphoproliferative disorder (PTLD) rates as well as mortality in pediatric liver transplant recipients.

Source

N.D. Crins, C. Roever, A.D. Goralczyk, T. Friede. [Interleukin-2 receptor antagonists for pediatric liver transplant recipients: A systematic review and meta-analysis of controlled studies](#). *Pediatric Transplantation*, **18**(8):839-850, 2014.

References

- T.G. Heffron et al. Pediatric liver transplantation with daclizumab induction therapy. *Transplantation*, **75**(12):2040-2043, 2003.
- N.E.M. Gibelli et al. [Basiliximab-chimeric anti-IL2-R monoclonal antibody in pediatric liver transplantation: comparative study](#). *Transplantation Proceedings*, **36**(4):956-957, 2004.
- S. Schuller et al. [Daclizumab induction therapy associated with tacrolimus-MMF has better outcome compared with tacrolimus-MMF alone in pediatric living donor liver transplantation](#). *Transplantation Proceedings*, **37**(2):1151-1152, 2005.
- R. Ganschow et al. [Long-term results of basiliximab induction immunosuppression in pediatric liver transplant recipients](#). *Pediatric Transplantation*, **9**(6):741-745, 2005.
- M. Spada et al. [Randomized trial of basiliximab induction versus steroid therapy in pediatric liver allograft recipients under tacrolimus immunosuppression](#). *American Journal of Transplantation*, **6**(8):1913-1921, 2006.
- J.M. Gras et al. [Steroid-free, tacrolimus-basiliximab immunosuppression in pediatric liver transplantation: Clinical and pharmacoeconomic study in 50 children](#). *Liver Transplantation*, **14**(4):469-477, 2008.

See Also

[GoralczykEtAl2011](#).

Examples

```
data("CrinsEtAl2014")
## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
crins.es <- escalc(measure="OR",
                  ai=exp.AR.events, n1i=exp.total,
                  ci=cont.AR.events, n2i=cont.total,
                  slab=publication, data=CrinsEtAl2014)

print(crins.es)

# analyze using weakly informative half-Cauchy prior for heterogeneity:
crins.ma <- bayesmeta(crins.es, tau.prior=function(t){dhalfcauchy(t,scale=1)})

# show results:
print(crins.ma)
forestplot(crins.ma)
plot(crins.ma)

# show heterogeneity posterior along with prior:
```

```

plot(crins.ma, which=4, prior=TRUE)

# perform meta analysis using 2 randomized studies only
# but use 4 non-randomized studies to inform heterogeneity prior:
crins.nrand <- bayesmeta(crins.es[crins.es$randomized=="no",],
                        tau.prior=function(t){dhalfcauchy(t,scale=1)})
crins.rand  <- bayesmeta(crins.es[crins.es$randomized=="yes",],
                        tau.prior=function(t){crins.nrand$dposterior(tau=t)})
plot(crins.nrand, which=4, prior=TRUE,
     main="non-randomized posterior = randomized prior")
plot(crins.rand, which=4, prior=TRUE, main="randomized posterior")
plot(crins.rand, which=1)

## End(Not run)

```

dhalflogistic	<i>Half-logistic distribution.</i>
---------------	------------------------------------

Description

Half-logistic density, distribution, quantile functions and random number generation.

Usage

```

dhalflogistic(x, scale=1, log=FALSE)
phalflogistic(q, scale=1)
qhalflogistic(p, scale=1)
rhalflogistic(n, scale=1)

```

Arguments

x, q	quantile.
p	probability.
n	number of observations.
scale	scale parameter (> 0).
log	logical; if TRUE, logarithmic density will be returned.

Details

The **half-logistic distribution** is simply a zero-mean logistic distribution that is restricted to take only positive values. If $X \sim \text{logistic}$, then $|sX| \sim \text{halflogistic}(\text{scale} = s)$.

Value

'dhalflogistic()' gives the density function, 'phalflogistic()' gives the cumulative distribution function (CDF), 'qhalflogistic()' gives the quantile function (inverse CDF), and 'rhalflogistic()' generates random deviates.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

N.L. Johnson, S. Kotz, N. Balakrishnan. *Continuous univariate distributions*, volume 2, chapter 23.11. Wiley, New York, 2nd edition, 1994.

See Also

[dlogis](#), [dhalfnormal](#), [dlomax](#), [drayleigh](#), [TurnerEtAlPrior](#), [RhodesEtAlPrior](#), [bayesmeta](#).

Examples

```
#####
# illustrate densities:
x <- seq(0,6,le=200)
plot(x, dhalfnormal(x), type="l", col="red", ylim=c(0,1),
      xlab=expression(tau), ylab=expression("probability density "*f(tau)))
lines(x, dhalflogistic(x), col="green3")
lines(x, dhalfcauchy(x), col="blue")
lines(x, dexp(x), col="cyan")
abline(h=0, v=0, col="grey")

# show log-densities (note the differing tail behaviour):
plot(x, dhalfnormal(x), type="l", col="red", ylim=c(0.001,1), log="y",
      xlab=expression(tau), ylab=expression("probability density "*f(tau)))
lines(x, dhalflogistic(x), col="green3")
lines(x, dhalfcauchy(x), col="blue")
lines(x, dexp(x), col="cyan")
abline(v=0, col="grey")
```

dhalfnormal

Half-normal, half-Student-t and half-Cauchy distributions.

Description

Half-normal, half-Student-t and half-Cauchy density, distribution, quantile functions and random number generation.

Usage

```
dhalfnormal(x, scale=1, log=FALSE)
phalfnormal(q, scale=1)
qhalfnormal(p, scale=1)
rhalfnormal(n, scale=1)

dhalft(x, scale=1, df, log=FALSE)
```

```

phalft(q, scale=1, df)
qhalft(p, scale=1, df)
rhalft(n, scale=1, df)

dhalfcauchy(x, scale=1, log=FALSE)
phalfcauchy(q, scale=1)
qhalfcauchy(p, scale=1)
rhalfcauchy(n, scale=1)

```

Arguments

x, q	quantile.
p	probability.
n	number of observations.
scale	scale parameter (> 0).
df	degrees-of-freedom parameter (> 0).
log	logical; if TRUE, logarithmic density will be returned.

Details

The **half-normal distribution** is simply a zero-mean normal distribution that is restricted to take only positive values. The *scale* parameter σ here corresponds to the underlying normal distribution's standard deviation: if $X \sim \text{Normal}(0, \sigma^2)$, then $|X| \sim \text{halfNormal}(\text{scale} = \sigma)$. Its mean is $\sigma\sqrt{2/\pi}$, and its variance is $\sigma^2(1 - 2/\pi)$. Analogously, the **half-t distribution** is a truncated Student-t distribution with *df* degrees-of-freedom, and the **half-Cauchy distribution** is again a special case of the half-t distribution with *df*=1 degrees of freedom.

Note that (half-) Student-t and Cauchy distributions arise as continuous *mixture distributions* of (half-) normal distributions. If

$$Y|\sigma \sim \text{Normal}(0, \sigma^2)$$

where the standard deviation is $\sigma = \sqrt{k/X}$ and X is drawn from a χ^2 -distribution with k degrees of freedom, then the marginal distribution of Y is Student-t with k degrees of freedom.

Value

'dhalfnormal()' gives the density function, 'phalfnormal()' gives the cumulative distribution function (CDF), 'qhalfnormal()' gives the quantile function (inverse CDF), and 'rhalfnormal()' generates random deviates. For the 'dhalft()', 'dhalfcauchy()' and related function it works analogously.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- A. Gelman. Prior distributions for variance parameters in hierarchical models. *Bayesian Analysis*, **1**(3):515-534, 2006.
- F. C. Leone, L. S. Nelson, R. B. Nottingham. **The folded normal distribution**. *Technometrics*, **3**(4):543-550, 1961.
- N. G. Polson, J. G. Scott. On the half-Cauchy prior for a global scale parameter. *Bayesian Analysis*, **7**(4):887-902, 2012.
- S. Psarakis, J. Panaretos. **The folded t distribution**. *Communications in Statistics - Theory and Methods*, **19**(7):2717-2734, 1990.

See Also

[dnorm](#), [dt](#), [dcauchy](#), [dlomax](#), [drayleigh](#), [TurnerEtAlPrior](#), [RhodesEtAlPrior](#), [bayesmeta](#).

Examples

```
#####
# illustrate densities:
x <- seq(0,6,le=200)
plot(x, dhalfnormal(x), type="l", col="red", ylim=c(0,1),
      xlab=expression(tau), ylab=expression("probability density " * f(tau)))
lines(x, dhalft(x, df=3), col="green")
lines(x, dhalfcauchy(x), col="blue")
lines(x, dexp(x), col="cyan")
abline(h=0, v=0, col="grey")

# show log-densities (note the differing tail behaviour):
plot(x, dhalfnormal(x), type="l", col="red", ylim=c(0.001,1), log="y",
      xlab=expression(tau), ylab=expression("probability density " * f(tau)))
lines(x, dhalft(x, df=3), col="green")
lines(x, dhalfcauchy(x), col="blue")
lines(x, dexp(x), col="cyan")
abline(v=0, col="grey")
```

dlomax

The Lomax distribution.

Description

Lomax density, distribution and quantile functions, and random number generation.

Usage

```
dlomax(x, scale=1, shape=1, log=FALSE)
plomax(q, scale=1, shape=1)
qlomax(p, scale=1, shape=1)
rlomax(n, scale=1, shape=1)
```

Arguments

x, q	quantile.
p	probability.
n	number of observations.
scale	scale parameter ($\lambda > 0$).
shape	shape parameter ($\alpha > 0$).
log	logical; if TRUE, logarithmic density will be returned.

Details

The Lomax distribution is a heavy-tailed distribution that also is a special case of a *Pareto distribution of the 2nd kind*. The probability density function of a Lomax distributed variable with shape $\alpha > 0$ and scale $\lambda > 0$ is given by

$$p(x) = (\alpha/\lambda)(1 + x/\lambda)^{-(\alpha+1)}.$$

The density function is monotonically decreasing in x . Its mean is $\lambda/(\alpha - 1)$ (for $\alpha > 1$) and its median is $\alpha(2^{1/\alpha} - 1)$. Its variance is finite only for $\alpha > 2$ and equals $(\lambda^2\alpha)/((\alpha - 1)^2(\alpha - 2))$. The cumulative distribution function (CDF) is given by

$$P(x) = 1 - (1 + x/\lambda)^{-\alpha}.$$

The Lomax distribution also arises as a **gamma-exponential mixture**. Suppose that X is a draw from an exponential distribution whose rate θ again is drawn from a gamma distribution with shape a and scale s (so that $E[\theta] = as$ and $\text{Var}(\theta) = as^2$, or $E[1/\theta] = \frac{1}{s(a+1)}$ and $\text{Var}(1/\theta) = \frac{1}{s^2(a-1)^2(a-2)}$). Then the marginal distribution of X is Lomax with scale $1/s$ and shape a . Consequently, if the moments of θ are given by $E[\theta] = \mu$ and $\text{Var}(\theta) = \sigma^2$, then X is Lomax distributed with shape $\alpha = (\frac{\mu}{\sigma})^2$ and scale $\lambda = \frac{\mu}{\sigma^2} = \frac{\alpha}{\mu}$. The gamma-exponential connection is also illustrated in an example below.

Value

'dlomax()' gives the density function, 'plomax()' gives the cumulative distribution function (CDF), 'qlomax()' gives the quantile function (inverse CDF), and 'rlomax()' generates random deviates.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

N.L. Johnson, S. Kotz, N. Balakrishnan. *Continuous univariate distributions*, volume 1. Wiley, New York, 2nd edition, 1994.

See Also

[dexp](#), [dgamma](#), [dhalfnormal](#), [dhalft](#), [dhalfcauchy](#), [drayleigh](#), [TurnerEtAlPrior](#), [RhodesEtAlPrior](#), [bayesmeta](#).

Examples

```
#####
# illustrate densities:
x <- seq(0,6,le=200)
plot(x, dexp(x, rate=1), type="l", col="cyan", ylim=c(0,1),
      xlab=expression(tau), ylab=expression("probability density " * f(tau)))
lines(x, dlomax(x), col="orange")
abline(h=0, v=0, col="grey")

# show log-densities (note the differing tail behaviour):
plot(x, dexp(x, rate=1), type="l", col="cyan", ylim=c(0.001,1), log="y",
      xlab=expression(tau), ylab=expression("probability density " * f(tau)))
lines(x, dlomax(x), col="orange")
abline(v=0, col="grey")

#####
# illustrate the gamma-exponential mixture connection;
# specify a number of samples:
N <- 10000
# specify some gamma shape and scale parameters
# (via mixing distribution's moments):
expectation <- 2.0
stdev <- 1.0
gammashape <- (expectation / stdev)^2
gammascall <- stdev^2 / expectation
print(c("expectation"=expectation, "stdev"=stdev,
        "shape"=gammashape, "scale"=gammascall))
# generate gamma-distributed rates:
lambda <- rgamma(N, shape=gammashape, scale=gammascall)
# generate exponential draws according to gamma-rates:
y <- rexp(N, rate=lambda)
# determine lomax quantiles accordingly parameterized:
x <- qlomax(ppoints(N), scale=1/gammascall, shape=gammashape)
# compare distributions in a Q-Q-plot:
plot(x, sort(y), log="xy", main="quantile-quantile plot",
      xlab="theoretical quantile", ylab="empirical quantile")
abline(0, 1, col="red")
```

drayleigh

The Rayleigh distribution.

Description

Rayleigh density, distribution, quantile function and random number generation.

Usage

```
drayleigh(x, scale=1, log=FALSE)
prayleigh(q, scale=1)
```

```
qrayleigh(p, scale=1)
rrayleigh(n, scale=1)
```

Arguments

x, q	quantile.
p	probability.
n	number of observations.
scale	scale parameter (> 0).
log	logical; if TRUE, logarithmic density will be returned.

Details

The Rayleigh distribution arises as the distribution of the square root of an exponentially distributed (or χ^2_2 -distributed) random variable. If X follows an exponential distribution with rate λ and expectation $1/\lambda$, then $Y = \sqrt{X}$ follows a Rayleigh distribution with scale $\sigma = 1/\sqrt{2\lambda}$ and expectation $\sqrt{\pi/(4\lambda)}$.

Note that the exponential distribution is the *maximum entropy distribution* among distributions supported on the positive real numbers and with a pre-specified expectation; so the Rayleigh distribution gives the corresponding distribution of its square root.

Value

'drayleigh()' gives the density function, 'prayleigh()' gives the cumulative distribution function (CDF), 'qrayleigh()' gives the quantile function (inverse CDF), and 'rrayleigh()' generates random deviates.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

N.L. Johnson, S. Kotz, N. Balakrishnan. *Continuous univariate distributions*, volume 1. Wiley, New York, 2nd edition, 1994.

See Also

[dexp](#), [dlomax](#), [dhalfnormal](#), [dhalft](#), [dhalfcauchy](#), [TurnerEtAlPrior](#), [RhodesEtAlPrior](#), [bayesmeta](#).

Examples

```
#####
# illustrate densities:
x <- seq(0,6,le=200)
plot(x, drayleigh(x, scale=0.5), type="l", col="green",
      xlab=expression(tau), ylab=expression("probability density "*f(tau)))
lines(x, drayleigh(x, scale=1/sqrt(2)), col="red")
lines(x, drayleigh(x, scale=1), col="blue")
```



```

abline(h=0, v=0, col="grey")

#####
# illustrate exponential / Rayleigh connection
# via a quantile-quantile plot (Q-Q-plot):
N <- 10000
exprate <- 5
plot(sort(sqrt(rexp(N, rate=exprate))),
      grayleigh(ppoints(N), scale=1/sqrt(2*exprate)))
abline(0, 1, col="red")

#####
# illustrate Maximum Entropy distributions
# under similar but different constraints:
mu <- 0.5
tau <- seq(0, 4*mu, le=100)
plot(tau, dexp(tau, rate=1/mu), type="l", col="red", ylim=c(0,1/mu),
      xlab=expression(tau), ylab="probability density")
lines(tau, drayleigh(tau, scale=1/sqrt(2*1/mu^2)), col="blue")
abline(h=0, v=0, col="grey")
abline(v=mu, col="darkgrey"); axis(3, at=mu, label=expression(mu))
# explicate constraints:
legend("topright", pch=15, col=c("red", "blue"),
      c(expression("Exponential: E[*tau*]"==mu),
        expression("Rayleigh: E[*tau^2*]"==mu^2)))

```

forest.bayesmeta	<i>Generate a forest plot for a bayesmeta object (based on the metafor package's plotting functions).</i>
------------------	---

Description

Generates a forest plot, showing individual estimates along with their 95 percent confidence intervals, resulting effect estimate and prediction interval.

Usage

```

## S3 method for class 'bayesmeta'
forest(x, xlab="effect size", refline=0, cex=1,...)

```

Arguments

x	a bayesmeta object.
xlab	title for the x-axis.
refline	value at which a vertical 'reference' line should be drawn (default is 0). The line can be suppressed by setting this argument to 'NA'.
cex	character and symbol expansion factor.
...	other arguments.

Details

Generates a simple forest plot illustrating the underlying data and resulting estimates (effect estimate and prediction interval).

Note

This function requires the **metafor** package to be installed.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

C. Lewis and M. Clarke. Forest plots: trying to see the wood and the trees. *BMJ*, **322**:1479, 2001.
 R.D. Riley, J.P. Higgins and J.J. Deeks. Interpretation of random effects meta-analyses. *BMJ*, **342**:d549, 2011.

See Also

[bayesmeta](#), [forest.default](#), [addpoly](#), [forestplot.bayesmeta](#)

Examples

```
data("CrinsEtAl2014")

## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
es.crins <- escalc(measure="OR",
                  ai=exp.AR.events, n1i=exp.total,
                  ci=cont.AR.events, n2i=cont.total,
                  slab=publication, data=CrinsEtAl2014)
# derive a prior distribution for the heterogeneity:
tp.crins <- TurnerEtAlPrior("surgical", "pharma", "placebo / control")
# perform meta-analysis:
ma.crins <- bayesmeta(es.crins, tau.prior=tp.crins$dprior)

#####
# plot:
forest(ma.crins, xlab="log odds ratio")

forest(ma.crins, trans=exp, refline=1, xlab="odds ratio")

## End(Not run)
```

forestplot.bayesmeta	<i>Generate a forest plot for a bayesmeta object (based on the forestplot package's plotting functions).</i>
----------------------	--

Description

Generates a forest plot, showing individual estimates along with their 95 percent confidence intervals, shrinkage intervals, resulting effect estimate and prediction interval.

Usage

```
## S3 method for class 'bayesmeta'
forestplot(x, labeltext, exponentiate=FALSE,
           prediction=TRUE, shrinkage=TRUE, digits=2, plot=TRUE,
           fn.ci_norm, fn.ci_sum, col, legend, boxsize, ...)
```

Arguments

x	a bayesmeta object.
labeltext	an (alternative) “labeltext” argument which is then handed on to the forestplot() function (see the help there). You can use this to change contents or add columns to the displayed table; see the example below.
exponentiate	a logical flag indicating whether to exponentiate numbers (effect sizes) in table and plot.
prediction	a logical flag indicating whether to show the prediction interval below the mean estimate.
shrinkage	a logical flag indicating whether to show shrinkage intervals along with the quoted estimates.
digits	The number of significant digits to be shown. This is interpreted relative to the standard errors of all estimates.
plot	a logical flag indicating whether to actually generate a plot.
fn.ci_norm, fn.ci_sum, col, legend, boxsize, ...	other arguments passed on to the forestplot package's forestplot function (see also the help there).

Details

Generates a forest plot illustrating the underlying data and resulting estimates (effect estimate and prediction interval, as well as shrinkage estimates and intervals).

Note

This function is based on the **forestplot** package's “[forestplot\(\)](#)” function.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. *arXiv preprint 1711.08683* (submitted for publication), 2017.
- C. Lewis and M. Clarke. Forest plots: trying to see the wood and the trees. *BMJ*, **322**:1479, 2001.
- C. Guddat, U. Grouven, R. Bender and G. Skipka. *A note on the graphical presentation of prediction intervals in random-effects meta-analyses*. *Systematic Reviews*, **1**(34), 2012.
- R.D. Riley, J.P. Higgins and J.J. Deeks. Interpretation of random effects meta-analyses. *BMJ*, **342**:d549, 2011.

See Also

[bayesmeta](#), [forestplot](#), [forest.bayesmeta](#), [plot.bayesmeta](#).

Examples

```
# load data:
data("CrinsEtAl2014")

## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
crins.es <- escalc(measure="OR",
                  ai=exp.AR.events, n1i=exp.total,
                  ci=cont.AR.events, n2i=cont.total,
                  slab=publication, data=CrinsEtAl2014)

print(crins.es)

# perform meta analysis:
crins.ma <- bayesmeta(crins.es, tau.prior=function(t){dhalfcauchy(t,scale=1)})

#####
# generate forest plots
require("forestplot")

# default options:
forestplot(crins.ma)

# exponentiate values (shown in table and plot), show vertical line at OR=1:
forestplot(crins.ma, expo=TRUE, zero=1)

# logarithmic x-axis:
forestplot(crins.ma, expo=TRUE, xlog=TRUE)

# omit prediction interval:
forestplot(crins.ma, predict=FALSE)
```

```

# omit shrinkage intervals:
forestplot(crins.ma, shrink=FALSE)

# show more decimal places:
forestplot(crins.ma, digits=3)

# change table values:
# (here: add columns for event counts)
fp <- forestplot(crins.ma, expo=TRUE, plot=FALSE)
labtext <- fp$labeltext
labtext <- cbind(labtext[,1],
                 c("treatment",
                   paste0(CrinsEtAl2014[, "exp.AR.events"], "/", CrinsEtAl2014[, "exp.total"]),
                   "", "")),
                 c("control",
                   paste0(CrinsEtAl2014[, "cont.AR.events"], "/", CrinsEtAl2014[, "cont.total"]),
                   "", "")),
                 labtext[,2:3])
labtext[1,4] <- "OR"
print(fp$labeltext) # before
print(labtext)      # after
forestplot(crins.ma, labeltext=labtext, expo=TRUE, xlog=TRUE)

# see also the "forestplot" help for more arguments that you may change,
# e.g. the "clip", "xticks", "xlab" and "title" arguments,
# or the "txt_gp" argument for label sizes etc.:
forestplot(crins.ma, clip=c(-4,1), xticks=(-3):0,
           xlab="log-OR", title="pediatric transplantation example",
           txt_gp = fpTxtGp(ticks = gpar(cex=1), xlab = gpar(cex=1)))

## End(Not run)

```

GoralczykEtAl2011

Liver transplant example data

Description

Numbers of cases (transplant patients) and events (acute rejections, steroid resistant rejections, and deaths) in experimental and control groups of 19 studies.

Usage

```
data("GoralczykEtAl2011")
```

Format

The data frame contains the following columns:

publication	character	publication identifier (first author and publication year)
--------------------	-----------	--

year	numeric	publication year
randomized	factor	randomization status (yes / no / not stated)
control.type	factor	type of control group ('concurrent' or 'historical')
comparison	factor	type of comparison ('IL-2RA only', 'delayed CNI', or 'no/low steroids')
IL2RA	factor	type of interleukin-2 receptor antagonist (IL-2RA) ('basiliximab' or 'daclizumab')
CNI	factor	type of calcineurin inhibitor (CNI) ('tacrolimus' or 'cyclosporine A')
MMF	factor	use of mycophenolate mofetil (MMF) (y/n)
followup	numeric	follow-up time in months
treat.AR.events	numeric	number of AR events in experimental group
treat.SRR.events	numeric	number of SRR events in experimental group
treat.deaths	numeric	number of deaths in experimental group
treat.total	numeric	number of cases in experimental group
control.AR.events	numeric	number of AR events in control group
control.SRR.events	numeric	number of SRR events in control group
control.deaths	numeric	number of deaths in control group
control.total	numeric	number of cases in control group

Details

A systematic literature review investigated the evidence on the effect of Interleukin-2 receptor antagonists (IL-2RA) and resulted in 19 controlled studies reporting acute rejection (AR) and steroid-resistant rejection (SRR) rates as well as mortality in adult liver transplant recipients.

Source

A.D. Goralczyk, N. Hauke, N. Bari, T.Y. Tsui, T. Lorf, A. Obed. **Interleukin-2 receptor antagonists for liver transplant recipients: A systematic review and meta-analysis of controlled studies.** *Hepatology*, **54**(2):541-554, 2011.

See Also

[CrinsEtAl2014.](#)

Examples

```
data("GoralczykEtAl2011")
## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
goralczyk.es <- escalc(measure="OR",
                      ai=exp.AR.events, n1i=exp.total,
                      ci=cont.AR.events, n2i=cont.total,
                      slab=publication, data=GoralczykEtAl2011)
print(goralczyk.es[,c(1,10,12,13,15,16,17)])

# analyze using weakly informative half-Cauchy prior for heterogeneity:
goralczyk.ma <- bayesmeta(goralczyk.es, tau.prior=function(t){dhalfcauchy(t,scale=1)})
```

```
# show summary:
print(goralczyk.ma)

# show forest plot:
forestplot(goralczyk.ma)

## End(Not run)
```

HinksEtAl2010

JIA example data

Description

Log odds ratios indicating association of a genetic variant (CCR5) with juvenile idiopathic arthritis (JIA).

Usage

```
data("HinksEtAl2010")
```

Format

The data frame contains the following columns:

study	character	publication identifier
year	numeric	publication year
country	character	country
or	numeric	odds ratio (OR)
or.lower	numeric	lower 95 percent confidence bound for OR
or.upper	numeric	upper 95 percent confidence bound for OR
log.or	numeric	logarithmic OR
log.or.se	numeric	standard error of logarithmic OR

Details

Results from a genetic association study (Hinks et al, 2010) were combined with data from two additional studies (Pralhad et al., 2006; Lindner et al., 2007) in order to determine the combined evidence regarding the association of a particular genetic marker (CCR5) with juvenile idiopathic arthritis (JIA).

Source

A. Hinks et al. [Association of the CCR5 gene with juvenile idiopathic arthritis](#). *Genes and Immunity*, 11(7):584-589, 2010.

References

- S. Prahalad et al. **Association of two functional polymorphisms in the CCR5 gene with juvenile rheumatoid arthritis** *Genes and Immunity*, 7:468-475, 2006.
- E. Lindner et al. **Lack of association between the chemokine receptor 5 polymorphism CCR5delta32 in rheumatoid arthritis and juvenile idiopathic arthritis.** *BMC Medical Genetics*, 8:33, 2007.
- C. Roever, G. Knapp, T. Friede. **Hartung-Knapp-Sidik-Jonkman approach and its modification for random-effects meta-analysis with few studies.** *BMC Medical Research Methodology*, 15:99, 2015.

Examples

```
data("HinksEtAl2010")

## Not run:
# perform meta analysis based on weakly informative half-normal prior:
bma01 <- bayesmeta(y      = HinksEtAl2010$log.or,
                  sigma  = HinksEtAl2010$log.or.se,
                  labels  = HinksEtAl2010$study,
                  tau.prior = function(t){dhalfnormal(t,scale=1.0)})

# perform meta analysis based on slightly more informative half-normal prior:
bma02 <- bayesmeta(y      = HinksEtAl2010$log.or,
                  sigma  = HinksEtAl2010$log.or.se,
                  labels  = HinksEtAl2010$study,
                  tau.prior = function(t){dhalfnormal(t,scale=0.5)})

# show heterogeneity posteriors:
par(mfrow=c(2,1))
plot(bma01, which=4, prior=TRUE, taulim=c(0,1))
plot(bma02, which=4, prior=TRUE, taulim=c(0,1))
par(mfrow=c(1,1))

# show heterogeneity estimates:
rbind("half-normal(1.0)"=bma01$summary[, "tau"],
      "half-normal(0.5)"=bma02$summary[, "tau"])
# show q-profile confidence interval for tau in comparison:
require("metafor")
ma03 <- rma.uni(yi=log.or, sei=log.or.se, slab=study, data=HinksEtAl2010)
confint(ma03)$random["tau",c("ci.lb","ci.ub")]
# show I2 values in the relevant range:
tau <- seq(0, 0.7, by=0.1)
cbind("tau"=tau,
      "I2" =bma01$I2(tau=tau))

# show effect estimates:
round(rbind("half-normal(1.0)" = bma01$summary[, "mu"],
           "half-normal(0.5)" = bma02$summary[, "mu"]), 5)

# show forest plot:
forestplot(bma02)
# show shrinkage estimates:
bma02$theta
```



```
## End(Not run)
```

normalmixture	<i>Compute normal mixtures</i>
---------------	--------------------------------

Description

This function allows to derive density, distribution function and quantile function of a normal mixture with fixed mean (μ) and random standard deviation (σ).

Usage

```
normalmixture(density,
               cdf = Vectorize(function(x){integrate(density,0,x)$value}),
               mu = 0, delta = 0.01, epsilon = 0.0001,
               rel.tol.integrate = 2^16*.Machine$double.eps,
               abs.tol.integrate = rel.tol.integrate,
               tol.uniroot = rel.tol.integrate)
```

Arguments

density	the σ mixing distribution's probability density function.
cdf	the σ mixing distribution's cumulative distribution function.
mu	the normal mean (μ).
delta, epsilon	the parameters specifying the desired accuracy for approximation of the mixing distribution, and with that determining the number of σ support points being used internally. Smaller values imply greater accuracy and greater computational burden (Roever and Friede, 2017).
rel.tol.integrate, abs.tol.integrate, tol.uniroot	the rel.tol, abs.tol and tol 'accuracy' arguments that are passed to the integrate() or uniroot() functions for internal numerical integration or root finding (see also the help there).

Details

When a normal random variable

$$X|\mu, \sigma \sim \text{Normal}(\mu, \sigma^2)$$

has a fixed mean μ , but a random standard deviation

$$\sigma|\phi \sim G(\phi)$$

following some probability distribution $G(\phi)$, then the *marginal distribution* of X ,

$$X|\mu, \phi$$

is a *mixture distribution* (Lindsay, 1995; Seidel, 2010).

The mixture distribution’s probability density function etc. result from integration and often are not available in analytical form. The `normalmixture()` function approximates density, distribution function and quantile function to some pre-set accuracy by a *discrete* mixture of normal distributions based on a finite number of σ values using the ‘DIRECT’ algorithm (Roever and Friede, 2017).

Either the “density” or “cdf” argument needs to be supplied. If only “density” is given, then the CDF is derived via integration, if only “cdf” is supplied, then the density function is not necessary.

In **meta-analysis** applications, mixture distributions arise e.g. in the context of **prior predictive distributions**. Assuming that a study-specific effect θ_i *a priori* is distributed as

$$\theta_i | \mu, \tau \sim \text{Normal}(\mu, \tau^2)$$

with a prior distribution for the heterogeneity τ ,

$$\tau | \phi \sim G(\phi)$$

yields a setup completely analogous to the above one.

Since it is sometimes hard to judge what constitutes a sensible heterogeneity prior, it is often useful to inspect the implications of certain settings in terms of the corresponding *prior predictive distribution* of

$$\theta_i | \mu, \phi$$

indicating the *a priori* implied variation between studies due to heterogeneity alone based on a certain prior setup (Spiegelhalter et al., 2004, Sec. 5.7.3). Some examples using different heterogeneity priors are illustrated below.

Value

A list containing the following elements:

delta, epsilon	The supplied design parameters.
mu	the normal mean.
bins	the number of bins used.
support	the matrix containing lower, upper and reference points for each bin and its associated probability.
density	the mixture’s density function(x).
cdf	the mixture’s cumulative distribution function(x) (CDF).
quantile	the mixture’s quantile function(p) (inverse CDF).
mixing.density	the mixing distribution’s density function() (if supplied).
mixing.cdf	the mixing distribution’s cumulative distribution function().

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- B.G. Lindsay. *Mixture models: theory, geometry and applications*. Vol. 5 of *CBMS Regional Conference Series in Probability and Statistics*, Institute of Mathematical Statistics, Hayward, CA, USA, 1995.
- C. Roever, T. Friede. [Discrete approximation of a mixture distribution via restricted divergence](#). *Journal of Computational and Graphical Statistics*, **26**(1):217-222, 2017.
- C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. [arXiv preprint 1711.08683](#) (submitted for publication), 2017.
- W.E. Seidel. Mixture models. In M. Lovric (ed.) *International Encyclopedia of Statistical Science*, Springer, Heidelberg, pp. 827-829, 2010.
- D.J. Spiegelhalter, K.R. Abrams, J.P. Myles. *Bayesian approaches to clinical trials and health-care evaluation*. Wiley & Sons, 2004.

See Also

[bayesmeta](#).

Examples

```
#####
# compare half-normal mixing distributions with different scales:
nm05 <- normalmixture(cdf=function(x){phalfnormal(x, scale=0.5)})
nm10 <- normalmixture(cdf=function(x){phalfnormal(x, scale=1.0)})
# (this corresponds to the case of assuming a half-normal prior
# for the heterogeneity tau)

# check the structure of the returned object:
str(nm05)

# show density functions:
# (these would be the marginal (prior predictive) distributions
# of study-specific effects theta[i])
x <- seq(-1, 3, by=0.01)
plot(x, nm05$density(x), type="l", col="blue", ylab="density")
lines(x, nm10$density(x), col="red")
abline(h=0, v=0, col="grey")

# show cumulative distributions:
plot(x, nm05$cdf(x), type="l", col="blue", ylab="CDF")
lines(x, nm10$cdf(x), col="red")
abline(h=0:1, v=0, col="grey")

# determine 5 percent and 95 percent quantiles:
rbind("HN(0.5)"=nm05$quantile(c(0.05,0.95)),
      "HN(1.0)"=nm10$quantile(c(0.05,0.95)))

#####
# compare different mixing distributions
# (half-normal, half-Cauchy, exponential and Lomax):
```

```

nmHN <- normalmixture(cdf=function(x){phalfnormal(x, scale=0.5)})
nmHC <- normalmixture(cdf=function(x){phalfcauchy(x, scale=0.5)})
nmE  <- normalmixture(cdf=function(x){pexp(x, rate=2)})
nmL  <- normalmixture(cdf=function(x){plomax(x, shape=4, scale=2)})

# show densities (logarithmic y-axis):
x <- seq(-1, 3, by=0.01)
plot(x, nmHN$density(x), col="green", type="l", ylab="density", ylim=c(0.005, 6.5), log="y")
lines(x, nmHC$density(x), col="red")
lines(x, nmE$density(x), col="blue")
lines(x, nmL$density(x), col="cyan")
abline(v=0, col="grey")

# show CDFs:
plot(x, nmHN$cdf(x), col="green", type="l", ylab="CDF", ylim=c(0,1))
lines(x, nmHC$cdf(x), col="red")
lines(x, nmE$cdf(x), col="blue")
lines(x, nmL$cdf(x), col="cyan")
abline(h=0:1, v=0, col="grey")
# add "exponential" x-axis at top:
axis(3, at=log(c(0.5,1,2,5,10,20)), lab=c(0.5,1,2,5,10,20))
# show 95 percent quantiles:
abline(h=0.95, col="grey", lty="dashed")
abline(v=nmHN$quantile(0.95), col="green", lty="dashed")
abline(v=nmHC$quantile(0.95), col="red", lty="dashed")
abline(v=nmE$quantile(0.95), col="blue", lty="dashed")
abline(v=nmL$quantile(0.95), col="cyan", lty="dashed")
rbind("half-normal(0.5)"=nmHN$quantile(0.95),
      "half-Cauchy(0.5)"=nmHC$quantile(0.95),
      "exponential(2.0)"=nmE$quantile(0.95),
      "Lomax(4,2)"      =nmL$quantile(0.95))

#####
# a normal mixture distribution example where the solution
# is actually known analytically: the Student-t distribution.
# If  $Y|\sigma \sim N(0, \sigma^2)$ , where  $\sigma = \sqrt{k/X}$ 
# and  $X|k \sim \text{Chi}^2(\text{df}=k)$ ,
# then the marginal  $Y|k$  is Student-t with k degrees of freedom.

# define CDF of sigma:
CDF <- function(sigma, df){pchisq(df/sigma^2, df=df, lower.tail=FALSE)}

# numerically approximate normal mixture (with k=5 d.f.):
k <- 5
nmT1 <- normalmixture(cdf=function(x){CDF(x, df=k)})
# in addition also try a more accurate approximation:
nmT2 <- normalmixture(cdf=function(x){CDF(x, df=k)}, delta=0.001, epsilon=0.00001)
# check: how many grid points were required?
nmT1$bins
nmT2$bins

# show true and approximate densities:

```

```

x <- seq(-2,10,le=400)
plot(x, dt(x, df=k), type="l")
abline(h=0, v=0, col="grey")
lines(x, nmT1$density(x), col="red", lty="dashed")
lines(x, nmT2$density(x), col="blue", lty="dotted")

# show ratios of true and approximate densities:
plot(x, nmT1$density(x)/dt(x, df=k), col="red",
     type="l", log="y", ylab="density ratio")
abline(h=1, v=0, col="grey")
lines(x, nmT2$density(x)/dt(x, df=k), col="blue")

```

Peto1980

*Aspirin after myocardial infarction example data***Description**

Numbers of cases (patients) and events (deaths) in treatment and control groups of six studies.

Usage

```
data("Peto1980")
```

Format

The data frame contains the following columns:

publication	character	publication identifier
treat.cases	numeric	number of cases in treatment group
treat.events	numeric	number of events in treatment group
control.cases	numeric	number of cases in control group
control.events	numeric	number of events in control group

Details

Quoting from Brockwell and Gordon (2001): “The collection consists of six studies, each examining the effect of aspirin after myocardial infarction. In each study the number of patients who died after having been given either aspirin or a control drug is recorded.”

Source

S.E. Brockwell, I.R. Gordon. [A comparison of statistical methods for meta-analysis](#). *Statistics in Medicine*, **20**(6):825-840, 2001.

References

R. Peto. [Aspirin after myocardial infarction](#). *The Lancet*, **315**(8179):1172-1173, 1980.

Examples

```
data("Peto1980")
## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
peto.es <- escalc(measure="OR",
                  ai=treat.events, n1i=treat.cases,
                  ci=control.events, n2i=control.cases,
                  slab=publication, data=Peto1980)

print(peto.es)

# check sensitivity to different prior choices:
peto.ma01 <- bayesmeta(peto.es)
peto.ma02 <- bayesmeta(peto.es, tau.prior=function(t){dhalfnormal(t, scale=1)})

par(mfrow=c(2,1))
plot(peto.ma01, which=4, prior=TRUE, taulim=c(0,1), main="uniform prior")
plot(peto.ma02, which=4, prior=TRUE, taulim=c(0,1), main="half-normal prior")
par(mfrow=c(1,1))

# compare heterogeneity (tau) estimates:
print(rbind("uniform"      =peto.ma01$summary[, "tau"],
            "half-normal" =peto.ma02$summary[, "tau"]))

# compare effect (mu) estimates:
print(rbind("uniform"      =peto.ma01$summary[, "mu"],
            "half-normal" =peto.ma02$summary[, "mu"]))

summary(peto.ma02)
forestplot(peto.ma02)
plot(peto.ma02)

## End(Not run)
```

plot.bayesmeta

Generate summary plots for a [bayesmeta](#) object.

Description

Generates a forest plot, and joint and marginal posterior density plots for the two parameters of the random-effects meta-analysis model.

Usage

```
## S3 method for class 'bayesmeta'
plot(x, main=deparse(substitute(x)),
     which=1:4, prior=FALSE, forest.margin=8,
     mulim=c(NA,NA), taulim=c(NA,NA),
     violin=FALSE, ...)
```

Arguments

<code>x</code>	a <code>bayesmeta</code> object.
<code>main</code>	a character string giving the main title for the plot(s).
<code>which</code>	an indicator of which plots to generate.
<code>prior</code>	an indicator whether to also draw the prior density in marginal posterior density plots.
<code>forest.margin</code>	the width of the margin to the left of the forest plot. This may require some manual tweaking so that the study labels fit properly.
<code>mulim, tauim</code>	(optional) ranges of effect (μ) and heterogeneity (τ) values to be used for plotting.
<code>violin</code>	an indicator whether to draw the forest plot as a “violin plot”.
<code>...</code>	other graphical parameters.

Details

Depending on the value of the `which` argument, one or several plots are generated, namely

1. a forest plot, including a 95% credible interval (diamond) and a 95% prediction interval (rectangle) for the effect μ . The shown intervals for μ are based on posterior medians and shortest credible intervals (from `x$summary`). If `violin=TRUE`, the forest plot is plotted as a “violin plot”, i.e., via Gaussian densities for the estimates y_i (and their associated uncertainties), and the posterior densities for the effect μ , and for the predictive distribution.
2. a plot of the joint posterior density of heterogeneity (τ) and effect (μ). Red lines trace the contours of constant density corresponding to approximate 2D credible regions (based on a χ^2 -approximation to the logarithmic posterior density) as labelled. The credible regions are only an approximation based on a ‘well-behaved’, unimodal posterior; contour lines are omitted if the posterior mode is not finite. Blue lines show the conditional mean effect μ as a function of the heterogeneity τ (solid line) along with conditional 95% confidence bounds (dashed lines). Green lines indicate marginal medians and shortest 95% credible intervals for τ and μ .
3. the marginal posterior probability density of the effect μ with median and shortest 95% credible interval indicated. Depending on the `prior` argument, a dashed line showing the prior density is added. Note that for improper priors the scaling is arbitrary and may be inappropriate for the plot.
4. the marginal posterior probability density of the heterogeneity τ with median and shortest 95% credible interval indicated. Depending on the `prior` argument, a dashed line showing the prior density is added. Note that for improper priors the scaling is arbitrary and may be inappropriate for the plot.

The joint posterior density plot (2) especially highlights the dependence of the effect estimate on the heterogeneity parameter. In a ‘conventional’ frequentist meta-analysis, one would commonly first estimate the heterogeneity τ , and then fix this value and estimate the effect μ based on the assumption that the heterogeneity estimate was the true value. In the joint density plot, this would correspond to considering vertical “slices” of the parameter space, a slice at $\tau = 0$ for the fixed-effects model, and a slice at a different τ value for the random-effects model, where the blue lines would then indicate the corresponding estimate and confidence interval for μ .

Note that when using the `prior=TRUE` argument, the added line may end up be outside the plotted range, especially when using improper priors with arbitrary normalisation (consider adding it “manually” instead).

Value

Returns the supplied bayesmeta object (x).

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. *arXiv preprint 1711.08683* (submitted for publication), 2017.
- C. Guddat, U. Grouven, R. Bender and G. Skipka. *A note on the graphical presentation of prediction intervals in random-effects meta-analyses*. *Systematic Reviews*, **1**(34), 2012.
- R.D. Riley, J.P. Higgins and J.J. Deeks. Interpretation of random effects meta-analyses. *BMJ*, **342**:d549, 2011.

See Also

[bayesmeta](#), [forestplot.bayesmeta](#)

Examples

```
## Not run:
# example data by Snedecor and Cochran:
data("SnedecorCochran")

# analyze using a weakly informative prior
# (may take a few seconds to compute!):
ma <- bayesmeta(y=SnedecorCochran[, "mean"], sigma=sqrt(SnedecorCochran[, "var"]),
               label=SnedecorCochran[, "no"],
               mu.prior.mean=50, mu.prior.sd=50,
               tau.prior=function(x){dhalfcauchy(x, scale=10)})

# show some plots:
plot(ma, main="Snedecor/Cochran data", prior=TRUE)
plot(ma, main="Snedecor/Cochran data", which=1, violin=TRUE)

## End(Not run)
```

pppvalue *Posterior predictive p-values*

Description

Compute posterior or prior predictive p -values from a [bayesmeta](#) object.

Usage

```
pppvalue(x, parameter = "mu", value = 0.0,
         alternative = c("two.sided", "less", "greater"),
         statistic = "median",
         rejection.region,
         n = 10,
         prior = FALSE,
         quietly = FALSE,
         parallel, seed, ...)
```

Arguments

<code>x</code>	a bayesmeta object.
<code>parameter</code>	the parameter to be tested. May be the effect ("mu"), the heterogeneity ("tau") or one of the study-specific (θ_i) parameters denoted by their label or their index.
<code>value</code>	the (null-) hypothesized value.
<code>alternative</code>	the type of alternative hypothesis.
<code>statistic</code>	the figure to be used as the 'test statistic', or 'discrepancy variable'. May be chosen as "t", "Q" or "cdf", or among the row names of the bayesmeta object's ' <code>...\$summary</code> ' element. <i>Or</i> it may be specified as a function. For details, see below.
<code>rejection.region</code>	the test statistic's rejection region. May be one of "upper.tail", "lower.tail" or "two.tailed". If unspecified, it is set automatically based on the ' <code>alternative</code> ' and ' <code>statistic</code> ' parameters.
<code>n</code>	the number of Monte Carlo replications to be generated. The default value is <code>n=10</code> , but in practice a substantially larger value should be appropriate.
<code>prior</code>	a logical flag to request <i>prior predictive</i> (instead of <i>posterior predictive</i>) p -values. Prior predictive values are only available for hypotheses concerning the effect (μ) and heterogeneity (τ) parameters.
<code>quietly</code>	a logical flag to show (or suppress) output during computation; this may also speed up computations slightly.
<code>parallel</code>	the number of parallel processes to utilize. By default, if multiple (k) cores are detected, then k-1 parallel processes are used.
<code>seed</code>	(optional) an integer random seed value to generate reproducible results.
<code>...</code>	further parameters passed to ' <code>statistic</code> ', <i>if</i> the ' <code>statistic</code> ' argument was specified as a function.

Details

Posterior predictive p -values are Bayesian analogues to ‘classical’ p -values (Meng, 1994; Gelman, Meng and Stern, 1996; Gelman et al., 2014). The `pppvalue()` function allows to compute these values for one- and two-sided hypotheses concerning the effect (μ) or heterogeneity (τ) parameter, or one of the study-specific effect parameters (θ_i) in a random-effects meta-analysis.

Prior predictive p -values have a similar interpretation, but they have a stronger dependence on the prior specification and are only available when the prior is proper; for a more detailed discussion, see Gelman, Meng and Stern (1996, Sec. 4).

The function may also be used to only generate samples $(\tau, \mu, \theta_i, y_i)$ without having to also derive a statistic or a p -value. In order to achieve that, the ‘`statistic`’ argument may be specified as ‘NA’, and generated samples may be recovered from the ‘`...$replicates`’ output element.

p -values from Monte Carlo sampling: The computation is based on Monte Carlo sampling and repeated analysis of re-generated data sets drawn from the parameters’ (conditional) posterior predictive (or prior) distribution; so the p -value derivation is somewhat computationally expensive. The p -value eventually is computed based on how far in the tail area the actual data are (in terms of the realized ‘test statistic’ or ‘discrepancy’) relative to the Monte-Carlo-sampled distribution. Accuracy of the computed p -value hence strongly depends on the number of samples (as specified through the ‘`n`’ argument) that are generated. Also, (near-) zero p -values need to be interpreted with caution, and in relation to the used Monte Carlo sample size (n).

‘Test’-statistics or ‘discrepancy variables’: The ‘`statistic`’ argument determines the statistic to be computed from the data as a measure of deviation from the null hypothesis. If specified as “`Q`”, then Cochran’s Q statistic is computed; this is useful for testing for homogeneity ($\tau = 0$). If specified as one of the row names of the ‘`x$summary`’ element, then, depending on the type of null hypothesis specified through the ‘`parameter`’ argument, the corresponding parameter’s posterior quantity is used for the statistic. If specified as “`t`”, then a t -type statistic is computed (the difference between the corresponding parameter’s posterior mean and its hypothesized value, divided by the posterior standard deviation). If specified as “`cdf`”, the parameter’s marginal posterior cumulative distribution function evaluated at the hypothesized value (‘`value`’) is used.

The ‘`statistic`’ argument may also be specified as an arbitrary function of the data (y). The function’s first argument then needs to be the data (y), additional arguments may be passed as arguments (‘`...`’) to the ‘`pppvalue()`’ function. See also the examples below.

One- and two-sided hypotheses: Specification of one- or two-sided hypotheses not only has implications for the determination of the p -value from the samples, but also for the sampling process itself. Parameter values are drawn from a subspace according to the null hypothesis, which for a two-sided test is a line, and for a one-sided test is a half-plane. This also implies that one- and two-sided p -values cannot simply be converted into one another.

For example, when specifying `pppvalue(..., param="mu", val=0, alt="two.sided")`, then first parameter values (τ, μ) are drawn from the conditional posterior distribution $p(\tau, \mu | y, \sigma, \mu = 0)$, and subsequently new data sets are generated based on the parameters. If a one-sided hypothesis is specified, e.g. via `pppvalue(..., param="mu", val=0, alt="less")`, then parameters are drawn from $p(\tau, \mu | y, \sigma, \mu > 0)$.

For a hypothesis concerning the individual effect parameters θ_i , conditions are imposed on the corresponding θ_i . For example, for a specification of `pppvalue(..., param=2, val=0, alt="less")`, the hypothesis concerns the $i=2$ nd study’s effect parameter θ_2 . First a sample is generated from

$p(\theta_2|y, \sigma, \theta_2 > 0)$. Then samples of μ and τ are generated by conditioning on the generated θ_2 value, and data y are generated by conditioning on all three.

Unless explicitly specified through the ‘rejection.region’ argument, the test statistic’s “rejection region” (the direction in which extreme statistic values indicate a departure from the null hypothesis) is set based on the ‘alternative’ and ‘statistic’ parameters. The eventually used setting can be checked in the output’s ‘...\$rejection.region’ component.

Computation: When aiming to compute a p -value, it is probably a good idea to first start with a smaller ‘n’ argument to get a rough idea of the p -value’s order of magnitude as well as the computational speed, before going over to a larger, more realistic n value. The implementation is able to utilize multiple processors or cores via the **parallel** package; details may be specified via the ‘parallel’ argument.

Value

A list of class ‘htest’ containing the following components:

statistic	the ‘test statistic’ (or ‘discrepancy’) value based on the actual data.
parameter	the number (n) of Monte Carlo replications used.
p.value	the derived p -value.
null.value	the (null-) hypothesized parameter value.
alternative	the type of alternative hypothesis.
method	a character string indicating what type of test was performed.
data.name	the name of the underlying bayesmeta object.
call	an object of class call giving the function call that generated the htest object.
rejection.region	the test statistic’s rejection region.
replicates	a list containing the replicated parameters (τ, μ, θ_i), data (y_i) and statistic, along with an indicator for those samples constituting the distribution’s ‘tail area’.
computation.time	The computation time (in seconds) used.

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

- X.-L. Meng. **Posterior predictive p-values**. *The Annals of Statistics*, **22**(3):1142-1160, 1994.
- A. Gelman, X.-L. Meng, H. Stern. Posterior predictive assessment of model fitness via realized discrepancies. *Statistica Sinica*, **6**(4):733-760, 1996.
- A. Gelman, J.B. Carlin, H.S. Stern, D.B. Dunson, A. Vehtari, D.B. Rubin. *Bayesian data analysis*. Chapman & Hall / CRC, Boca Raton, 2014.
- C. Roever. Bayesian random-effects meta-analysis using the bayesmeta R package. *arXiv preprint 1711.08683* (submitted for publication), 2017.

See Also

[bayesmeta](#), [prop.test](#).

Examples

```
## Not run:
# perform a meta analysis;
# load data:
data("CrinsEtAl2014")
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
require("metafor")
crins.srr <- escalc(measure="OR",
                   ai=exp.SRR.events, n1i=exp.total,
                   ci=cont.SRR.events, n2i=cont.total,
                   slab=publication, data=CrinsEtAl2014, subset=c(1,4,6))

# analyze:
bma <- bayesmeta(crins.srr, mu.prior.mean=0, mu.prior.sd=4,
                 tau.prior=function(t){dhalfnormal(t, scale=0.5)})

# compute a 2-sided p-value for the effect (mu) parameter
# (note: this may take a while!):
p <- pppvalue(bma, parameter="mu", value=0, n=100)

# show result:
print(p)

# show the test statistic's distribution
# along with its actualized value:
plot(ecdf(p$replicates$statistic[,1]),
     xlim=range(c(p$statistic, p$replicates$statistic[,1])))
abline(v=p$statistic, col="red")

# show the parameter values
# drawn from the (conditional) posterior distribution:
plot(bma, which=2)
abline(h=p$null.value) # (the null-hypothesized mu value)
points(p$replicates$tau, p$replicates$mu, col="cyan") # (the samples)

#####
# Among the 3 studies, only the first (Heffron, 2003) was randomized.
# One might wonder about this particular study's effect (theta[1])
# in the light of the additional evidence and compute a one-sided
# p-value:

p <- pppvalue(bma, parameter="Heffron", value=0, n=100, alternative="less")
print(p)

#####
# One may also define one's own 'test' statistic to be used.
# For example, one could utilize the Bayes factor to generate
# a p-value for the homogeneity (tau=0) hypothesis:
```

```

BF <- function(y, sigma)
{
  bm <- bayesmeta(y=y, sigma=sigma,
                 mu.prior.mean=0, mu.prior.sd=4,
                 tau.prior=function(t){dhalfnormal(t, scale=0.5)},
                 interval.type="central")
  # (central intervals are faster to compute;
  # interval type otherwise is not relevant here)
  return(bm$bayesfactor[1,"tau=0"])
}
# NOTE: the 'bayesmeta()' arguments above should probably match
#       the specifications from the original analysis

p <- pppvalue(bma, parameter="tau", statistic=BF, value=0, n=100,
             alternative="greater", rejection.region="lower.tail",
             sigma=bma$sigma)

print(p)

#####
# If one is only interested in generating samples (and not in test
# statistics or p-values), one may specify the 'statistic' argument
# as 'NA'.
# Note that different 'parameter', 'value' and 'alternative' settings
# imply different sampling schemes.

p <- pppvalue(bma, parameter="mu", statistic=NA, value=0,
             alternative="less", n=100)

plot(bma, which=2)
abline(h=p$null.value)
points(p$replicates$tau, p$replicates$mu, col="cyan")

## End(Not run)

```

RhodesEtAlPrior

Heterogeneity priors for continuous outcomes (standardized mean differences) as proposed by Rhodes et al. (2015).

Description

Use the prior specifications proposed in the paper by Rhodes et al., based on an analysis of studies using standardized mean differences (SMD) that were published in the *Cochrane Database of Systematic Reviews*.

Usage

```

RhodesEtAlPrior(outcome=c(NA, "obstetric outcome",
                          "resource use and hospital stay / process",
                          "internal and external structure-related outcome",

```

```
"general physical health and adverse event and pain and quality of life / functioning",
paste("signs / symptoms reflecting continuation / end of condition and infection",
      "/ onset of new acute / chronic disease"),
"mental health outcome", "biological marker", "various subjectively measured outcomes"),
comparator1=c("pharmacological", "non-pharmacological", "placebo / control"),
comparator2=c("pharmacological", "non-pharmacological", "placebo / control"),
area=c("other", "respiratory", "cancer"))
```

Arguments

outcome	The type of outcome investigated (see below for a list of possible values). The default (NA) is the general (marginal) setting, without considering meta-analysis characteristics as covariates.
comparator1	One comparator's type.
comparator2	The other comparator's type.
area	The medical area.

Details

Rhodes et al. conducted an analysis of studies listed in the *Cochrane Database of Systematic Reviews* that were investigating standardized mean differences (SMD) as endpoints. As a result, they proposed empirically motivated log-Student- t prior distributions for the (squared!) heterogeneity parameter τ^2 , depending on the particular type of outcome investigated and the type of comparison in question. The underlying t -distribution's location and scale parameters here are internally stored in a 3-dimensional array (named `RhodesEtAlParameters`) and are most conveniently accessed using the `RhodesEtAlPrior()` function.

The outcome argument specifies the type of outcome investigated. It may take one of the following values (partial matching is supported):

- NA
- "obstetric outcomes"
- "resource use and hospital stay / process"
- "internal and external structure-related outcome"
- "general physical health and adverse event and pain and quality of life / functioning"
- "signs / symptoms reflecting continuation / end of condition and infection / onset of new acute /
- "mental health outcome"
- "biological marker"
- "various subjectively measured outcomes".

Specifying "outcome=NA" (the default) yields the *marginal* setting, without considering meta-analysis characteristics as covariates.

The comparator1 and comparator2 arguments together specify the type of comparison in question. These may take one of the following values (partial matching is supported):

- "pharmacological"
- "non-pharmacological"

- "placebo / control".

Any combination is allowed for the `comparator1` and `comparator2` arguments, as long as not both arguments are set to "placebo / control". The `area` argument specifies the medical context; possible values are:

- "respiratory"
- "cancer"
- "other" (the default).

Note that the location and scale parameters refer to the logarithmic (*squared*) heterogeneity parameter τ^2 , which is modelled using a Student-*t* distribution with 5 degrees of freedom. When you want to use the prior specifications for τ , the square root, as the parameter (as is necessary when using the `bayesmeta()` function), you need to correct for the square root transformation. Taking the square root is equivalent to dividing by two on the log-scale, so the square root will still be log-Student-*t* distributed, but with halved location and scale parameters. The relevant transformations are already taken care of when using the resulting `$dprior()`, `$pprior()` and `$qprior()` functions; see also the example below.

Value

a list with elements

<code>parameters</code>	the location and scale parameters (corresponding to the logarithmic <i>squared</i> heterogeneity parameter τ^2 as well as τ).
<code>outcome.type</code>	the corresponding type of outcome.
<code>comparison.type</code>	the corresponding type of comparison.
<code>medical.area</code>	the medical context.
<code>dprior</code>	a <code>function(tau)</code> returning the prior density of τ .
<code>pprior</code>	a <code>function(tau)</code> returning the prior cumulative distribution function (CDF) of τ .
<code>qprior</code>	a <code>function(p)</code> returning the prior quantile function (inverse CDF) of τ .

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

K.M. Rhodes, R.M. Turner, J.P.T. Higgins. Predictive distributions were developed for the extent of heterogeneity in meta-analyses of continuous outcome data. *Journal of Clinical Epidemiology*, 68(1):52-60, 2015.

See Also

[TurnerEtAlPrior](#).

Examples

```

# determine prior distribution for a specific setting:
RP <- RhodesEtAlPrior("obstetric", "pharma", "placebo")
print(RP$parameters)
str(RP)
# a prior 95 percent interval for tau:
RP$qprior(c(0.025,0.975))

# the general (marginal) setting:
RP <- RhodesEtAlPrior()
print(RP$parameters)
str(RP)
# a prior 95 percent interval for tau:
RP$qprior(c(0.025,0.975))

## Not run:
# load "metafor" package:
require("metafor")
# load data:
data("dat.normand1999")
# compute effect sizes (standardized mean differences):
es <- escalc(measure="SMD", m1i=m1i, sd1i=sd1i, n1i=n1i,
             m2i=m2i, sd2i=sd2i, n2i=n2i,
             slab=source, data=dat.normand1999)

# derive appropriate prior:
RP <- RhodesEtAlPrior("resource use", "non-pharma", "non-pharma")
# show (central) prior 95 percent interval:
RP$qprior(c(0.025, 0.975))
# show prior 95 percent upper limit:
RP$qprior(0.95)

# perform meta analysis:
bma <- bayesmeta(es, tau.prior=RP$dprior)
# show results:
print(bma)
plot(bma, which=4, prior=TRUE)

## End(Not run)

```

Rubin1981

8-schools example data

Description

SAT coaching experiments in 8 schools.

Usage

```
data("Rubin1981")
```


Format

The data frame contains the following columns:

school	character	school identifier
effect	numeric	effect estimate
stderr	numeric	associated standard error

Details

Quoting from Gelman et al. (1997), Sec. 5.5: “A study was performed for the Educational Testing Service to analyze the effects of special coaching programs for SAT-V (Scholastic Aptitude Test-Verbal) in each of eight high schools. The outcome variable in each study was the score on a special administration of the SAT-V, a standardized multiple choice test administered by the Educational Testing Service and used to help colleges make admissions decisions; the scores can vary between 200 and 800, with mean about 500 and standard deviation about 100. The SAT examinations are designed to be resistant to short-term efforts directed specifically toward improving performance on the test; instead they are designed to reflect knowledge acquired and abilities developed over many years of education. Nevertheless, each of the eight schools in this study considered its short-term coaching program to be very successful at increasing SAT scores. Also, there was no prior reason to believe that any of the eight programs was more effective than any other or that some were more similar in effect to each other than to any other.”

Source

A. Gelman, J.B. Carlin, H. Stern, and D.B. Rubin. *Bayesian data analysis*. Chapman & Hall / CRC, Boca Raton, 1997.

References

D.B. Rubin. Estimation in parallel randomized experiments. *Journal of Educational Statistics*, 6(4):377-401, 1981.

A. Gelman. Prior distributions for variance parameters in hierarchical models. *Bayesian Analysis*, 1(3):515-534, 2006.

Examples

```
data("Rubin1981")

## Not run:
# analysis using weakly informative half-Cauchy prior
# (may take a few seconds to compute!):
schools <- bayesmeta(y=Rubin1981[, "effect"], sigma=Rubin1981[, "stderr"],
                    labels=Rubin1981[, "school"],
                    tau.prior=function(x){return(dhalfcauchy(x, scale=25))})

# show summary:
summary(schools)

# show shrinkage effect for 8 individual estimates:
schools$theta
```

```
## End(Not run)
```

SidikJonkman2007

Postoperative complication odds example data

Description

This data set contains the outcomes from 29 randomized clinical trials comparing the odds of post-operative complications in laparoscopic inguinal hernia repair (LIHR) versus conventional open inguinal hernia repair (OIHR).

Usage

```
data("SidikJonkman2007")
```

Format

The data frame contains the following columns:

id	character	identifier used in original publication by Memon et al. (2003)
id.sj	numeric	identifier used by Sidik and Jonkman (2007)
year	numeric	publication year
lihr.events	numeric	number of events under LIHR
lihr.cases	numeric	number of cases under LIHR
oihr.events	numeric	number of events under OIHR
oihr.cases	numeric	number of cases under OIHR

Details

Analysis may be done based on the logarithmic odds ratios:

$$\log(\text{lihr.events}) - \log(\text{lihr.cases} - \text{lihr.events}) - \log(\text{oihr.events}) + \log(\text{oihr.cases} - \text{oihr.events})$$

and corresponding standard errors:

$$\sqrt{1/\text{lihr.events} + 1/(\text{lihr.cases} - \text{lihr.events}) + 1/\text{oihr.events} + 1/(\text{oihr.cases} - \text{oihr.events})}$$

(you may also leave these computations to the **metafor** package's `escalc()` function).

The data set was used to compare different estimators for the (squared) heterogeneity τ^2 . The values yielded for this data set were (see Tab.1 in Sidik and Jonkman (2007)):

method of moments (MM)	0.429
variance component (VC)	0.841
maximum likelihood (ML)	0.562
restricted ML (REML)	0.598
empirical Bayes (EB)	0.703
model error variance (MV)	0.818
variation of MV (MVvc)	0.747

Source

M.A. Memon, N.J. Cooper, B. Memon, M.I. Memon, and K.R. Abrams. [Meta-analysis of randomized clinical trials comparing open and laparoscopic inguinal hernia repair](#). *British Journal of Surgery*, **90**(12):1479-1492, 2003.

References

K. Sidik and J.N. Jonkman. [A comparison of heterogeneity variance estimators in combining results of studies](#). *Statistics in Medicine*, **26**(9):1964-1981, 2007.

Examples

```
data("SidikJonkman2007")
# add log-odds-ratios and corresponding standard errors:
sj <- SidikJonkman2007
sj <- cbind(sj, "log.or"=log(sj[, "lihr.events"])-log(sj[, "lihr.cases"]-sj[, "lihr.events"]))
      -log(sj[, "oihr.events"])+log(sj[, "oihr.cases"]-sj[, "oihr.events"]),
      "log.or.se"=sqrt(1/sj[, "lihr.events"] + 1/(sj[, "lihr.cases"]-sj[, "lihr.events"])
      + 1/sj[, "oihr.events"] + 1/(sj[, "oihr.cases"]-sj[, "oihr.events"])))

## Not run:
# analysis using weakly informative Cauchy prior
# (may take a few seconds to compute!):
ma <- bayesmeta(y=sj[, "log.or"], sigma=sj[, "log.or.se"], label=sj[, "id.sj"],
  tau.prior=function(t){dhalfcauchy(t, scale=1)})

# show heterogeneity's posterior density:
plot(ma, which=4, main="Sidik/Jonkman example", prior=TRUE)

# show some numbers (mode, median and mean):
abline(v=ma$summary[c("mode", "median", "mean"), "tau"], col="blue")

# compare with Sidik and Jonkman's estimates:
sj.estimates <- sqrt(c("MM" = 0.429, # method of moments estimator
  "VC" = 0.841, # variance component type estimator
  "ML" = 0.562, # maximum likelihood estimator
  "REML" = 0.598, # restricted maximum likelihood estimator
  "EB" = 0.703, # empirical Bayes estimator
  "MV" = 0.818, # model error variance estimator
  "MVvc" = 0.747)) # a variation of the MV estimator
abline(v=sj.estimates, col="red", lty="dashed")

# generate forest plot:
fp <- forestplot(ma, exponentiate=TRUE, plot=FALSE)
# add extra columns for ID and year:
labtext <- fp$labeltext
labtext[1,1] <- "ID 2"
labtext[31:32,1] <- ""
labtext <- cbind(c("ID 1", SidikJonkman2007[, "id"], "mean", "prediction"),
  labtext[,1],
  c("year", as.character(SidikJonkman2007[, "year"]), "", ""),
  labtext[, -1])
```

```
# plot:
forestplot(ma, labeltext=labtext, exponentiate=TRUE,
           xlog=TRUE, xlab="odds ratio", xticks=c(0.1,1,10))

## End(Not run)
```

SnedecorCochran

Artificial insemination of cows example data

Description

This data set gives means and (squared) standard errors of percentages of conceptions obtained from samples for six bulls.

Usage

```
data("SnedecorCochran")
```

Format

The data frame contains the following columns:

no	character	identifier
n	numeric	sample size
mean	numeric	mean
var	numeric	variance (squared standard error)

Details

Quoting from Snedecor and Cochran (1967), Sec. 10.18: “In research on artificial insemination of cows, a series of semen samples from a bull are sent out and tested for their ability to produce conceptions. The following data from a larger set kindly supplied by Dr. G. W. Salisbury, show the percentages of conceptions obtained from the samples for six bulls.”

Source

J. Hartung, G. Knapp, and B.K. Sinha. *Statistical meta-analysis with applications*. Wiley, Hoboken, NJ, USA, 2008.

References

G.W. Snedecor and W.G. Cochran. *Statistical Methods*. Iowa State University Press, Ames, IA, USA, 6th edition, 1967.

Examples

```
data("SnedecorCochran")
```

```
## Not run:
# analyze using uniform prior:
bma1 <- bayesmeta(y=SnedecorCochran[, "mean"],
                  sigma=sqrt(SnedecorCochran[, "var"]),
                  label=SnedecorCochran[, "no"],
                  tau.prior="uniform")

# analyze using Jeffreys prior:
bma2 <- bayesmeta(y=SnedecorCochran[, "mean"],
                  sigma=sqrt(SnedecorCochran[, "var"]),
                  label=SnedecorCochran[, "no"],
                  tau.prior="Jeffreys")

# compare results:
print(bma1)
print(bma2)

forestplot(bma1)
forestplot(bma2)

## End(Not run)
```

TurnerEtAlPrior	<i>(Log-Normal) heterogeneity priors for binary outcomes as proposed by Turner et al. (2015).</i>
-----------------	---

Description

Use the prior specifications proposed in the paper by Turner et al., based on an analysis of studies using binary endpoints that were published in the *Cochrane Database of Systematic Reviews*.

Usage

```
TurnerEtAlPrior(outcome=c(NA, "all-cause mortality", "obstetric outcomes",
                           "cause-specific mortality / major morbidity event / composite (mortality or morbidity)",
                           "resource use / hospital stay / process", "surgical / device related success / failure",
                           "withdrawals / drop-outs", "internal / structure-related outcomes",
                           "general physical health indicators", "adverse events",
                           "infection / onset of new disease",
                           "signs / symptoms reflecting continuation / end of condition", "pain",
                           "quality of life / functioning (dichotomized)", "mental health indicators",
                           "biological markers (dichotomized)", "subjective outcomes (various)"),
                 comparator1=c("pharmacological", "non-pharmacological", "placebo / control"),
                 comparator2=c("pharmacological", "non-pharmacological", "placebo / control"))
```

Arguments

outcome	The type of outcome investigated (see below for a list of possible values).
comparator1	One comparator's type.
comparator2	The other comparator's type.

Details

Turner et al. conducted an analysis of studies listed in the *Cochrane Database of Systematic Reviews* that were investigating binary endpoints. As a result, they proposed empirically motivated log-normal prior distributions for the (squared!) heterogeneity parameter τ^2 , depending on the particular type of outcome investigated and the type of comparison in question. The log-normal parameters (μ and σ) here are internally stored in a 3-dimensional array (named `TurnerEtAlParameters`) and are most conveniently accessed using the `TurnerEtAlPrior()` function.

The `outcome` argument specifies the type of outcome investigated. It may take one of the following values (partial matching is supported):

- NA
- "all-cause mortality"
- "obstetric outcomes"
- "cause-specific mortality / major morbidity event / composite (mortality or morbidity)"
- "resource use / hospital stay / process"
- "surgical / device related success / failure"
- "withdrawals / drop-outs"
- "internal / structure-related outcomes"
- "general physical health indicators"
- "adverse events"
- "infection / onset of new disease"
- "signs / symptoms reflecting continuation / end of condition"
- "pain"
- "quality of life / functioning (dichotomized)"
- "mental health indicators"
- "biological markers (dichotomized)"
- "subjective outcomes (various)"

Specifying "`outcome=NA`" (the default) yields the *marginal* setting, without considering meta-analysis characteristics as covariates.

The `comparator1` and `comparator2` arguments together specify the type of comparison in question. These may take one of the following values (partial matching is supported):

- "pharmacological"
- "non-pharmacological"
- "placebo / control"

Any combination is allowed for the `comparator1` and `comparator2` arguments, as long as not both arguments are set to "placebo / control".

Note that the log-normal prior parameters refer to the (*squared*) heterogeneity parameter τ^2 . When you want to use the prior specifications for τ , the square root, as the parameter (as is necessary when using the `bayesmeta()` function), you need to correct for the square root transformation. Taking the square root is equivalent to dividing by two on the log-scale, so the square root's distribution will still be log-normal, but with halved mean and standard deviation. The relevant transformations are already taken care of when using the resulting `$dprior()`, `$pprior()` and `$qprior()` functions; see also the example below.

Value

a list with elements

<code>parameters</code>	the log-normal parameters (μ and σ , corresponding to the <i>squared</i> heterogeneity parameter τ^2 as well as τ).
<code>outcome.type</code>	the corresponding type of outcome.
<code>comparison.type</code>	the corresponding type of comparison.
<code>dprior</code>	a function(<code>tau</code>) returning the prior density of τ .
<code>pprior</code>	a function(<code>tau</code>) returning the prior cumulative distribution function (CDF) of τ .
<code>qprior</code>	a function(<code>p</code>) returning the prior quantile function (inverse CDF) of τ .

Author(s)

Christian Roever <christian.roever@med.uni-goettingen.de>

References

R.M. Turner, D. Jackson, Y. Wei, S.G. Thompson, J.P.T. Higgins. [Predictive distributions for between-study heterogeneity and simple methods for their application in Bayesian meta-analysis](#). *Statistics in Medicine*, **34**(6):984-998, 2015.

See Also

[dlnorm](#), [RhodesEtAlPrior](#).

Examples

```
# load example data:
data("CrinsEtAl2014")

# determine corresponding prior parameters:
TP <- TurnerEtAlPrior("surgical", "pharma", "placebo / control")
print(TP)
# a prior 95 percent interval for tau:
TP$qprior(c(0.025, 0.975))
```

```

## Not run:
# compute effect sizes (log odds ratios) from count data
# (using "metafor" package's "escalc()" function):
crins.es <- escalc(measure="OR",
                  ai=exp.AR.events, n1i=exp.total,
                  ci=cont.AR.events, n2i=cont.total,
                  slab=publication, data=CrinsEtAl2014)
print(crins.es)

# perform meta analysis:
crins.ma01 <- bayesmeta(crins.es, tau.prior=TP$dprior)
# for comparison perform analysis using weakly informative Cauchy prior:
crins.ma02 <- bayesmeta(crins.es, tau.prior=function(t){dhalfcauchy(t,scale=1)})

# show results:
print(crins.ma01)
print(crins.ma02)
# compare estimates; heterogeneity (tau):
rbind("Turner prior"=crins.ma01$summary["tau"], "Cauchy prior"=crins.ma02$summary["tau"])
# effect (mu):
rbind("Turner prior"=crins.ma01$summary["mu"], "Cauchy prior"=crins.ma02$summary["mu"])

# illustrate heterogeneity priors and posteriors:
par(mfcol=c(2,2))
plot(crins.ma01, which=4, prior=TRUE, taulim=c(0,2),
     main="informative log-normal prior")
plot(crins.ma02, which=4, prior=TRUE, taulim=c(0,2),
     main="weakly informative half-Cauchy prior")
plot(crins.ma01, which=3, mulim=c(-3,0),
     main="informative log-normal prior")
abline(v=0, lty=3)
plot(crins.ma02, which=3, mulim=c(-3,0),
     main="weakly informative half-Cauchy prior")
abline(v=0, lty=3)
par(mfrow=c(1,1))

# compare prior and posterior 95 percent upper limits for tau:
TP$qprior(0.95)
crins.ma01$qposterior(0.95)
qhalfcauchy(0.95)
crins.ma02$qposterior(0.95)

## End(Not run)

```


qhalfcauchy (dhalfnormal), [19](#)
qhalflogistic (dhalflogistic), [18](#)
qhalfnormal (dhalfnormal), [19](#)
qhalfnormal (dhalfnormal), [19](#)
qhalfnormal (dhalfnormal), [19](#)
qlomax (dlomax), [21](#)
qrayleigh (drayleigh), [23](#)

rhalfcauchy (dhalfnormal), [19](#)
rhalflogistic (dhalflogistic), [18](#)
rhalfnormal (dhalfnormal), [19](#)
rhalfnormal (dhalfnormal), [19](#)
RhodesEtAlParameters (RhodesEtAlPrior),
[45](#)
RhodesEtAlPrior, [7](#), [19](#), [21](#), [22](#), [24](#), [45](#), [55](#)
rlomax (dlomax), [21](#)
rrayleigh (drayleigh), [23](#)
Rubin1981, [48](#)

SidikJonkman2007, [50](#)
SnedecorCochran, [52](#)
summary.bayesmeta (bayesmeta), [3](#)

TurnerEtAlParameters (TurnerEtAlPrior),
[53](#)
TurnerEtAlPrior, [7](#), [19](#), [21](#), [22](#), [24](#), [47](#), [53](#)

uniroot, [4](#), [10](#), [33](#)